

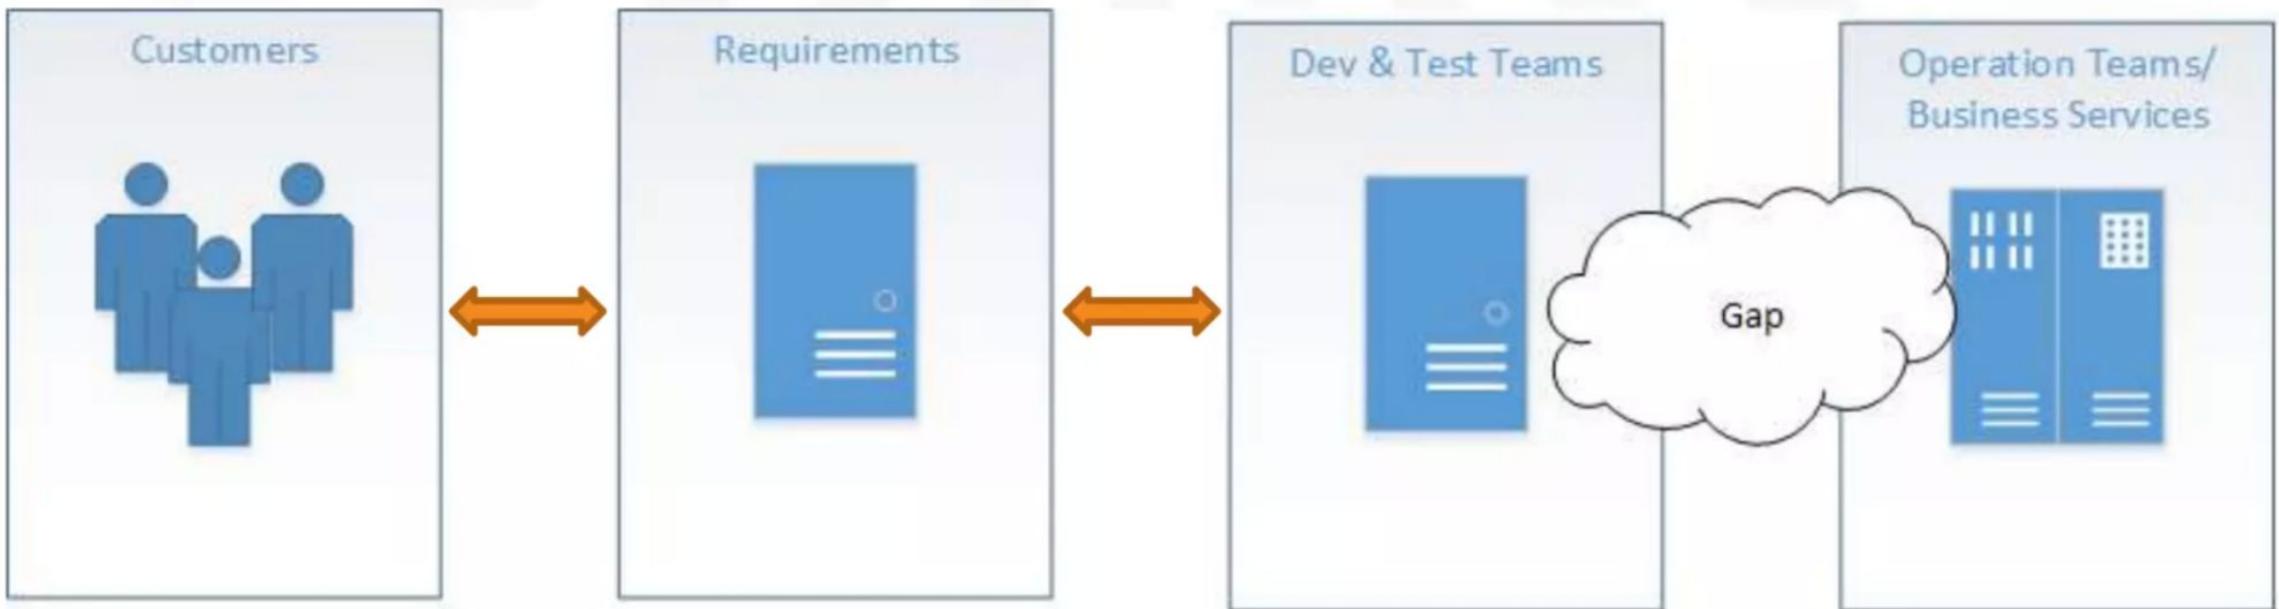
Kubernetes

Before DevOps

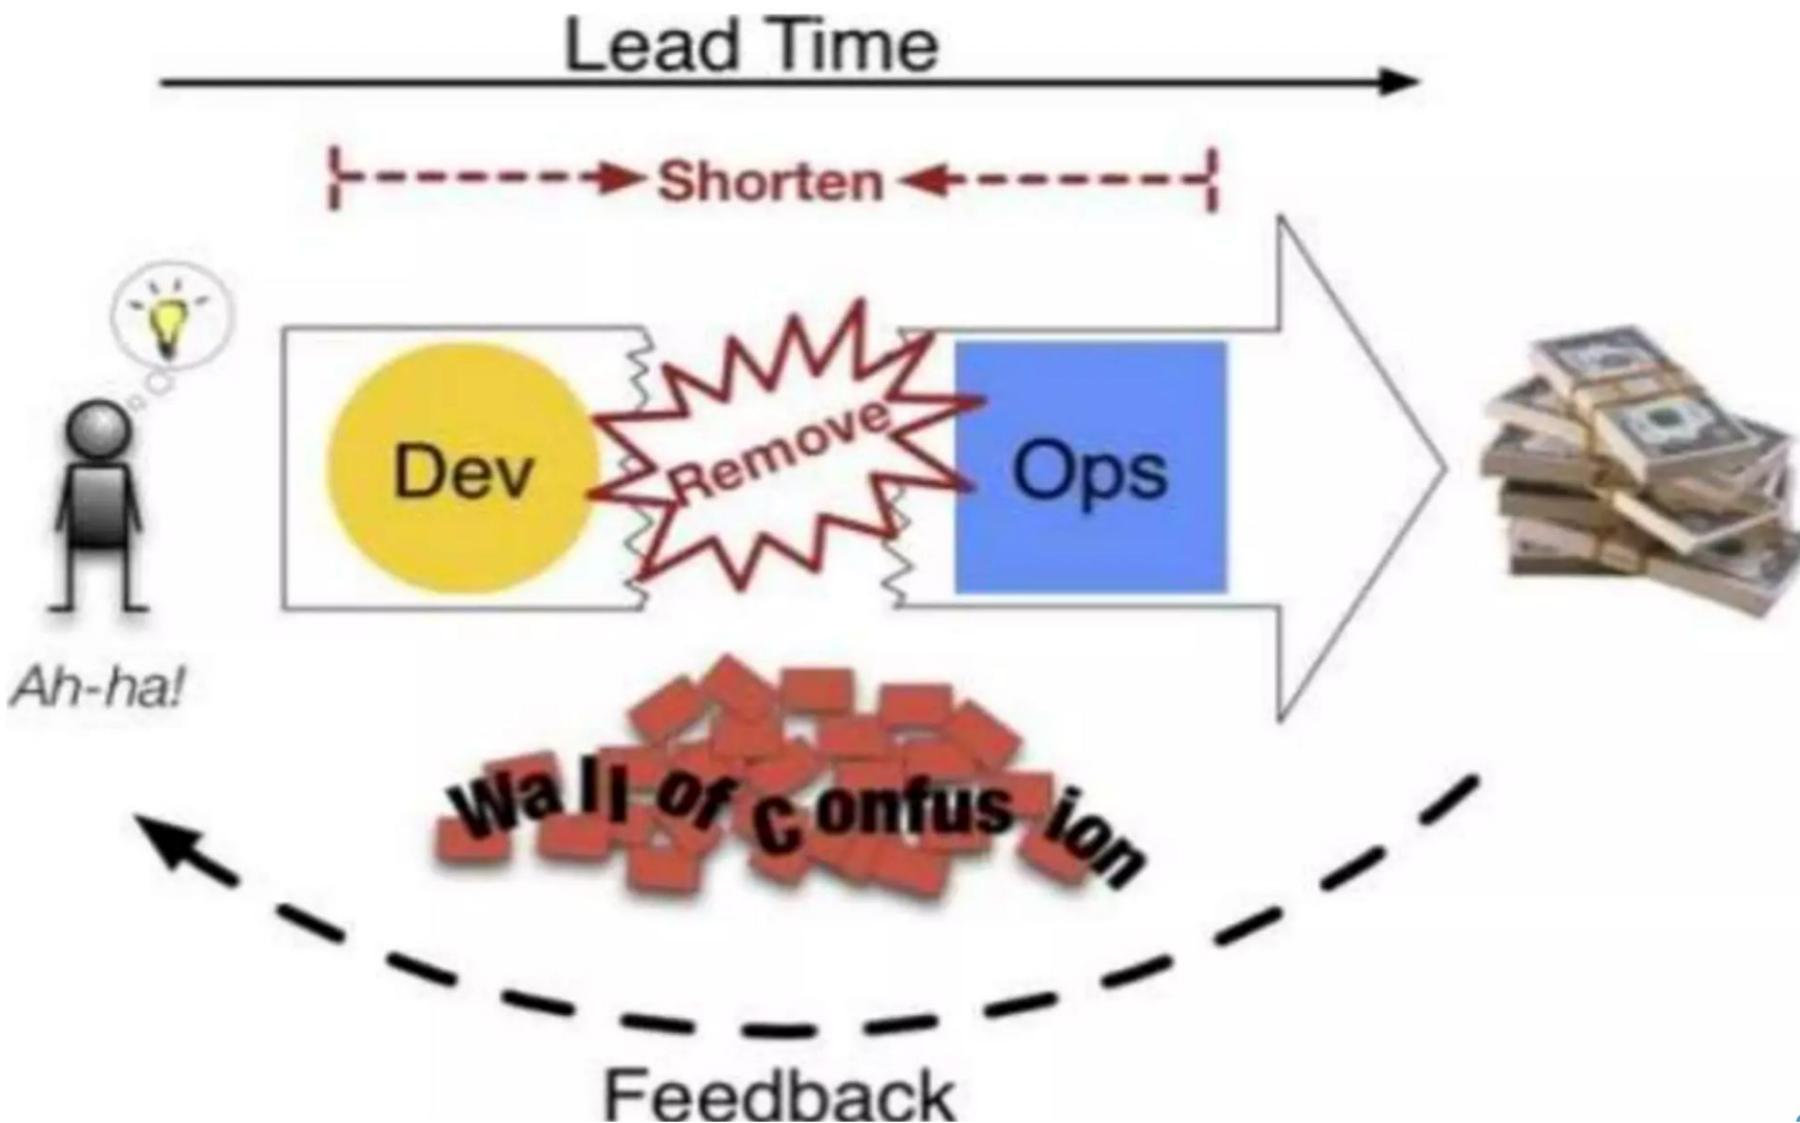




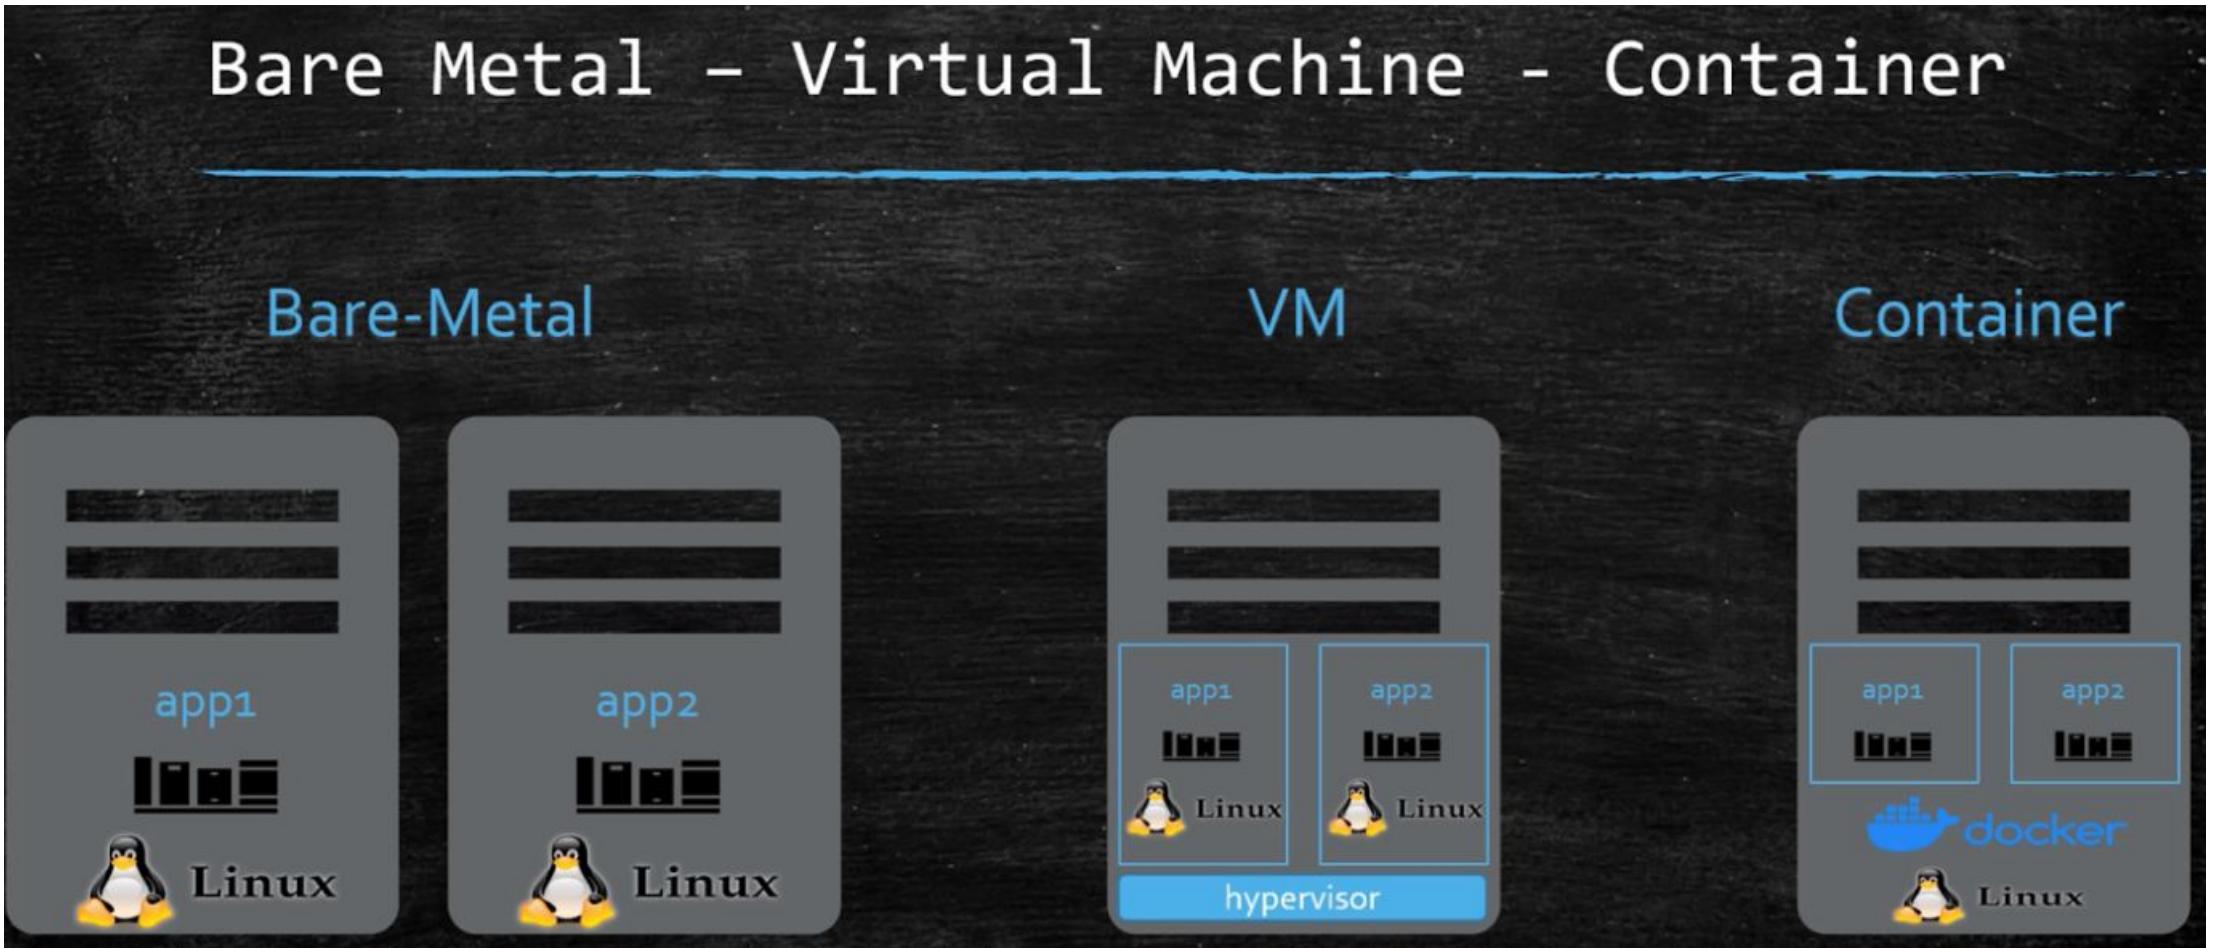
Why DevOps?



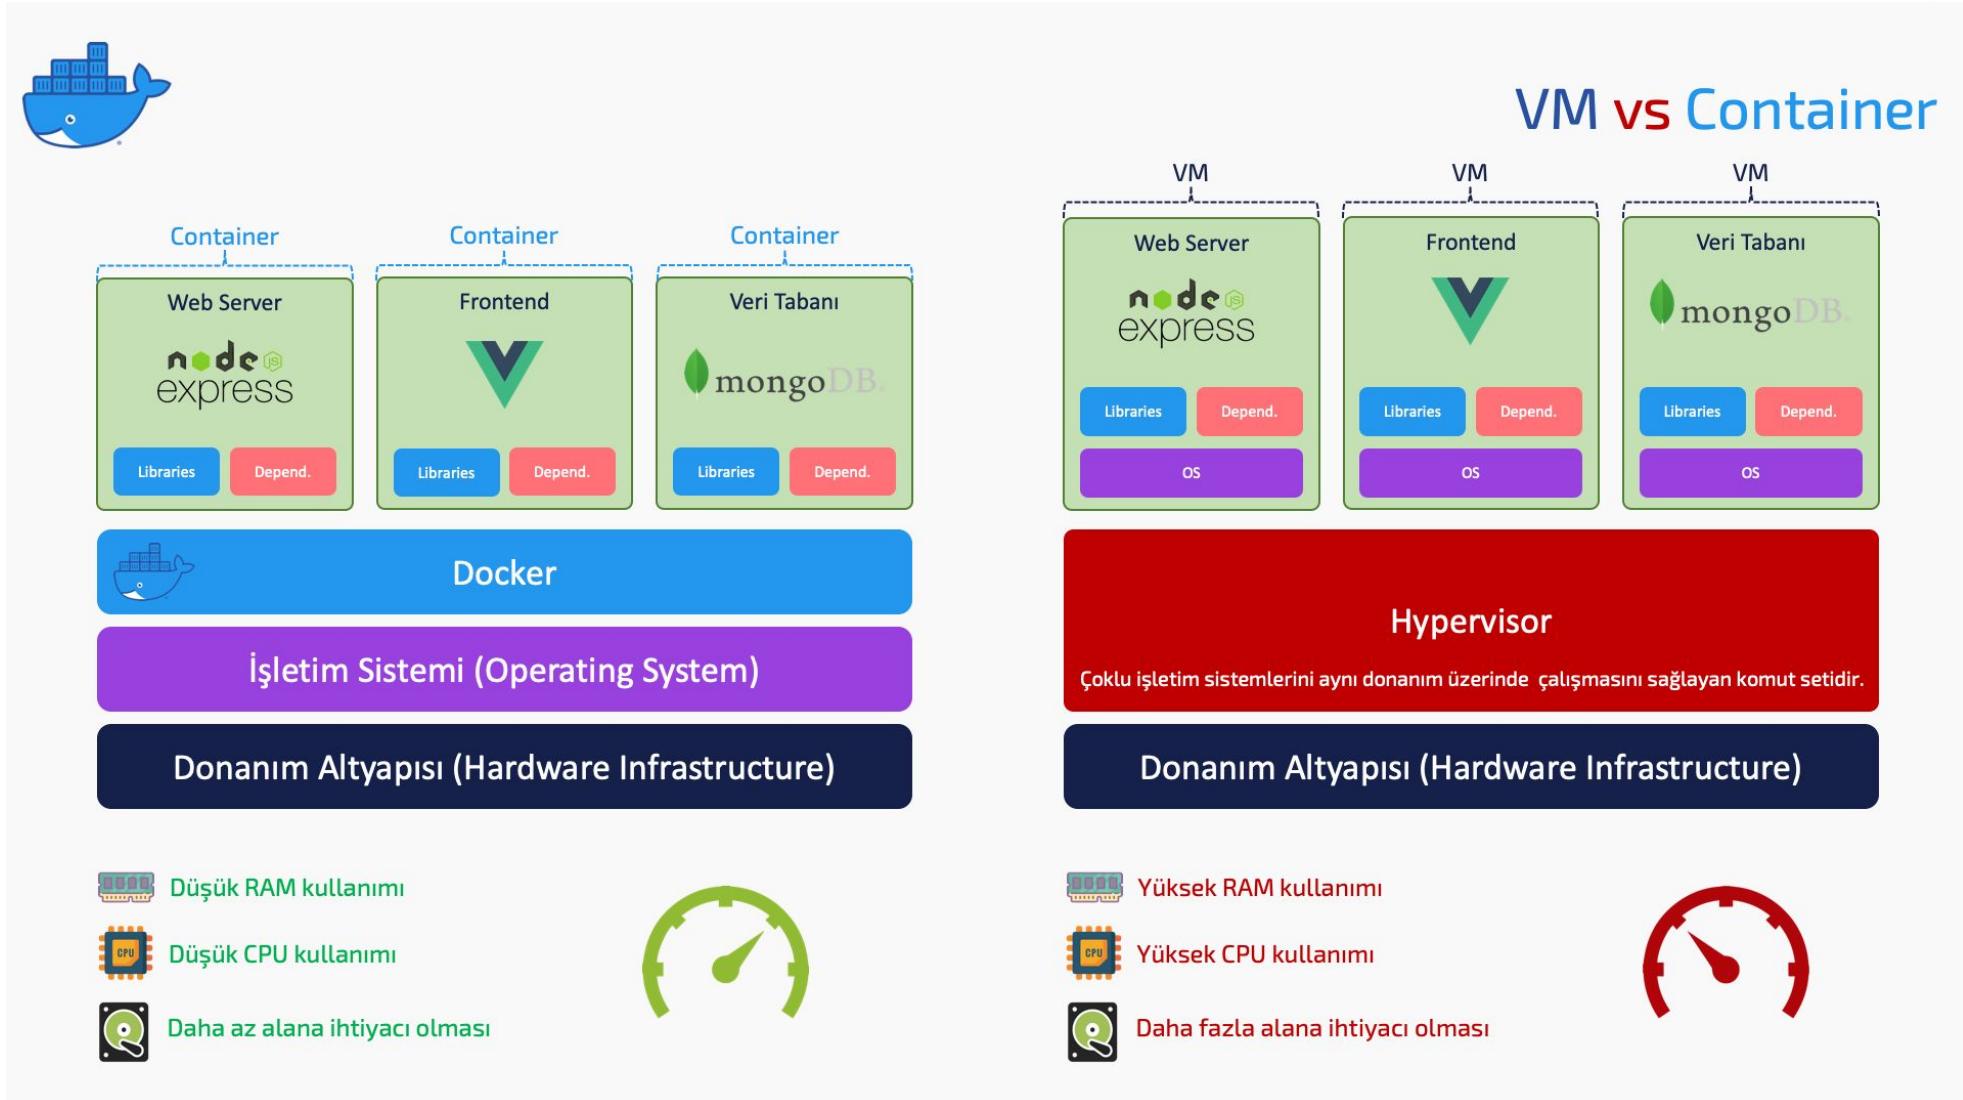
Why DevOps?



Containerization



Containerization

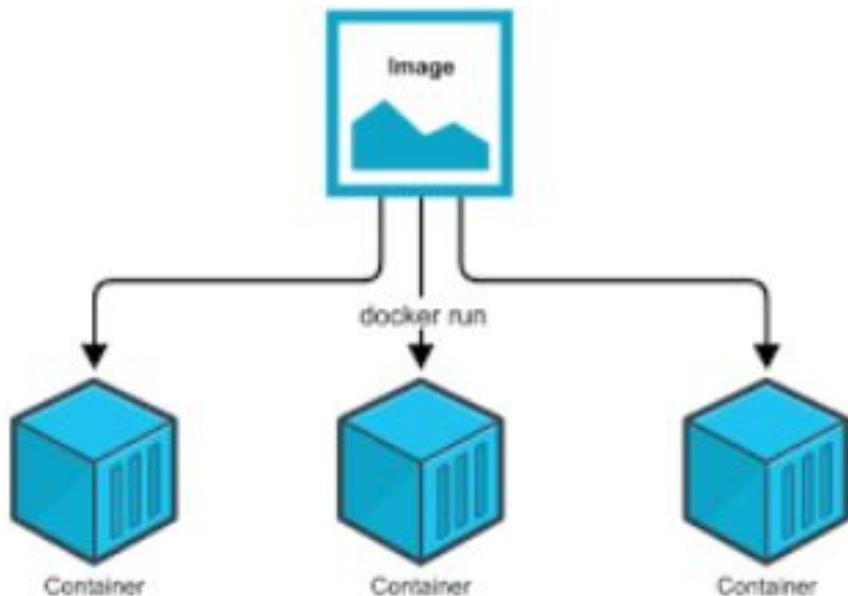
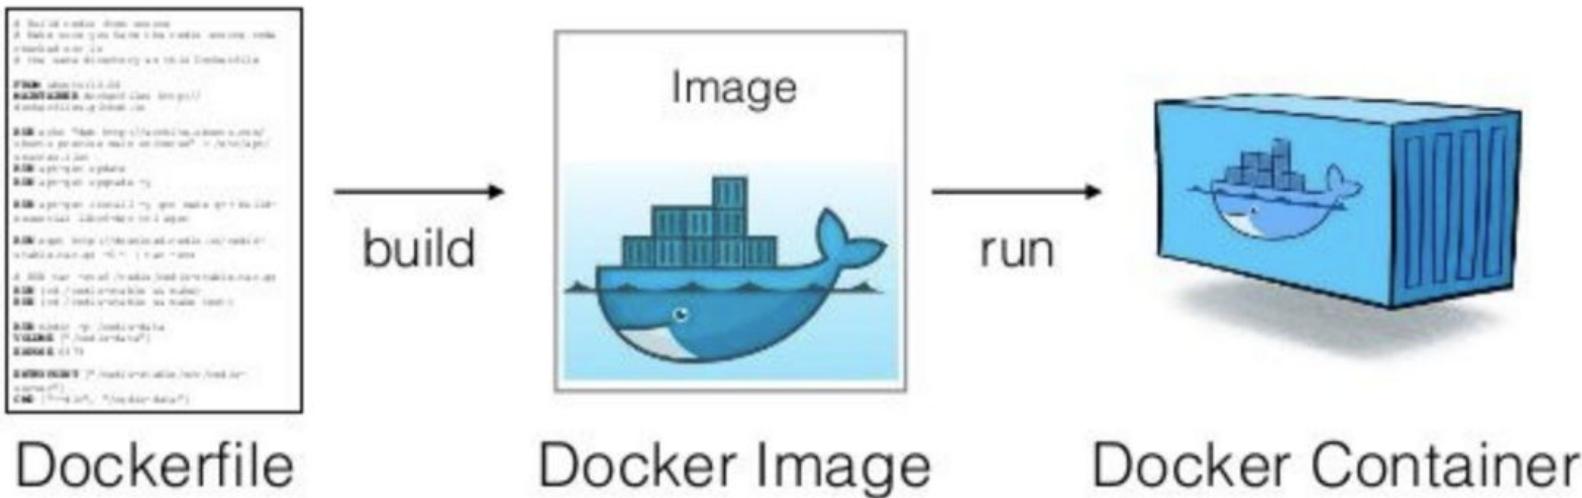


Containerization

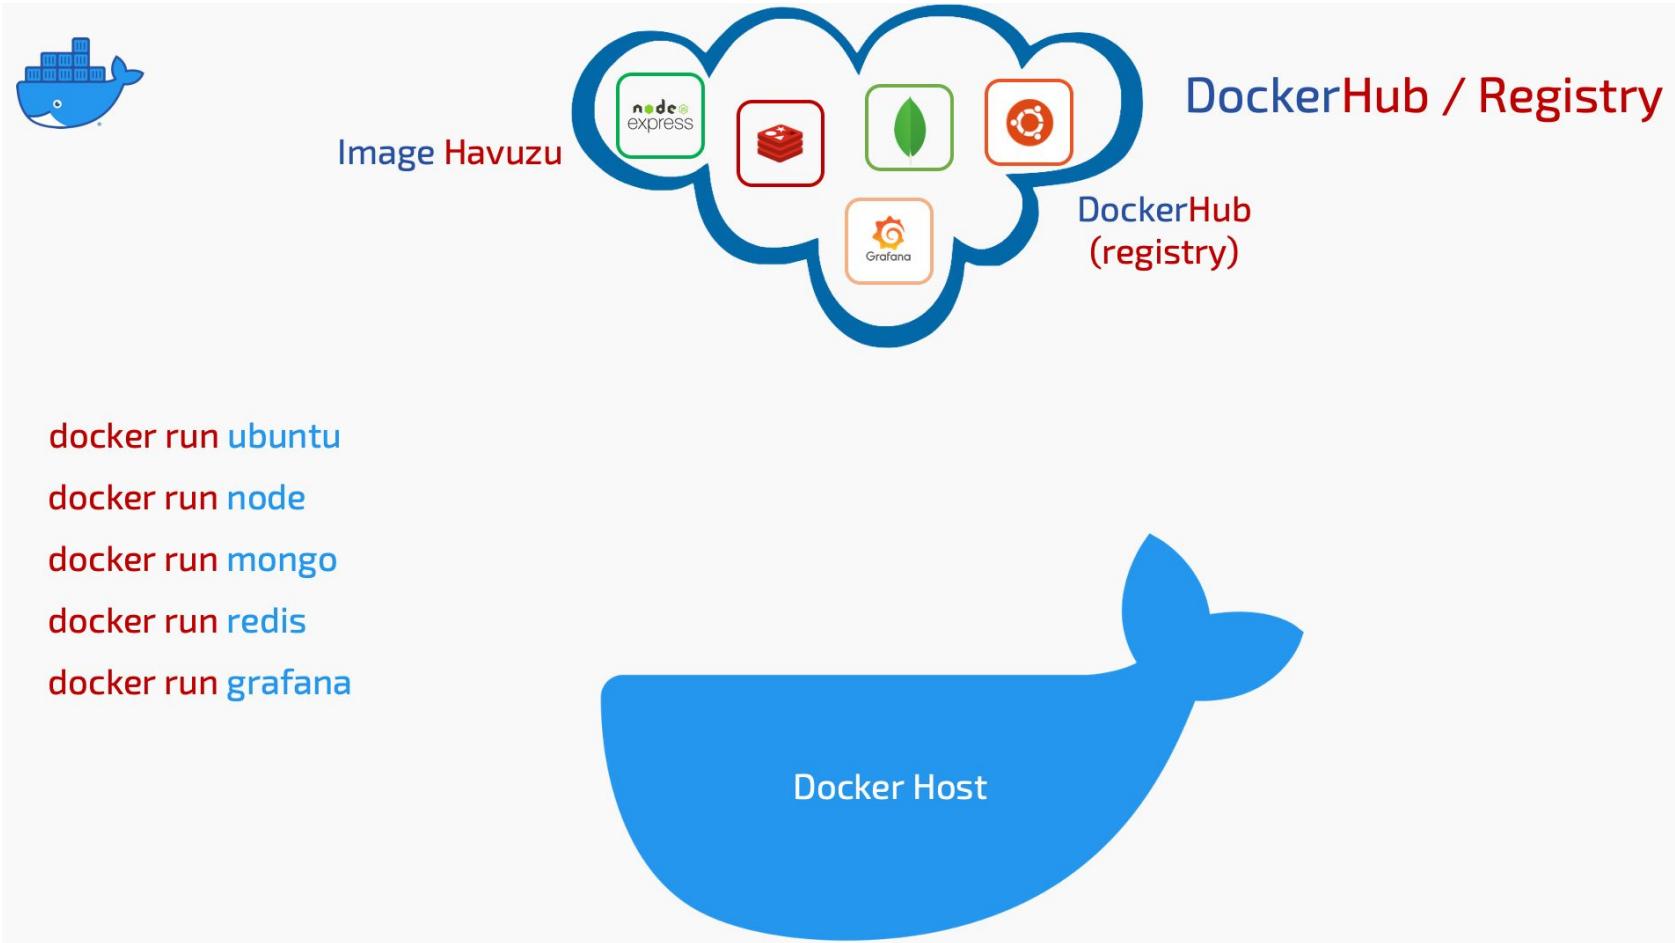
-  Application Isolation
-  Rapid Deployment and Scalability
-  Portability
-  Lightweight and Speed
-  Continuous Integration and Delivery (CI/CD)
-  Orchestration (Kubernetes)



Containerization

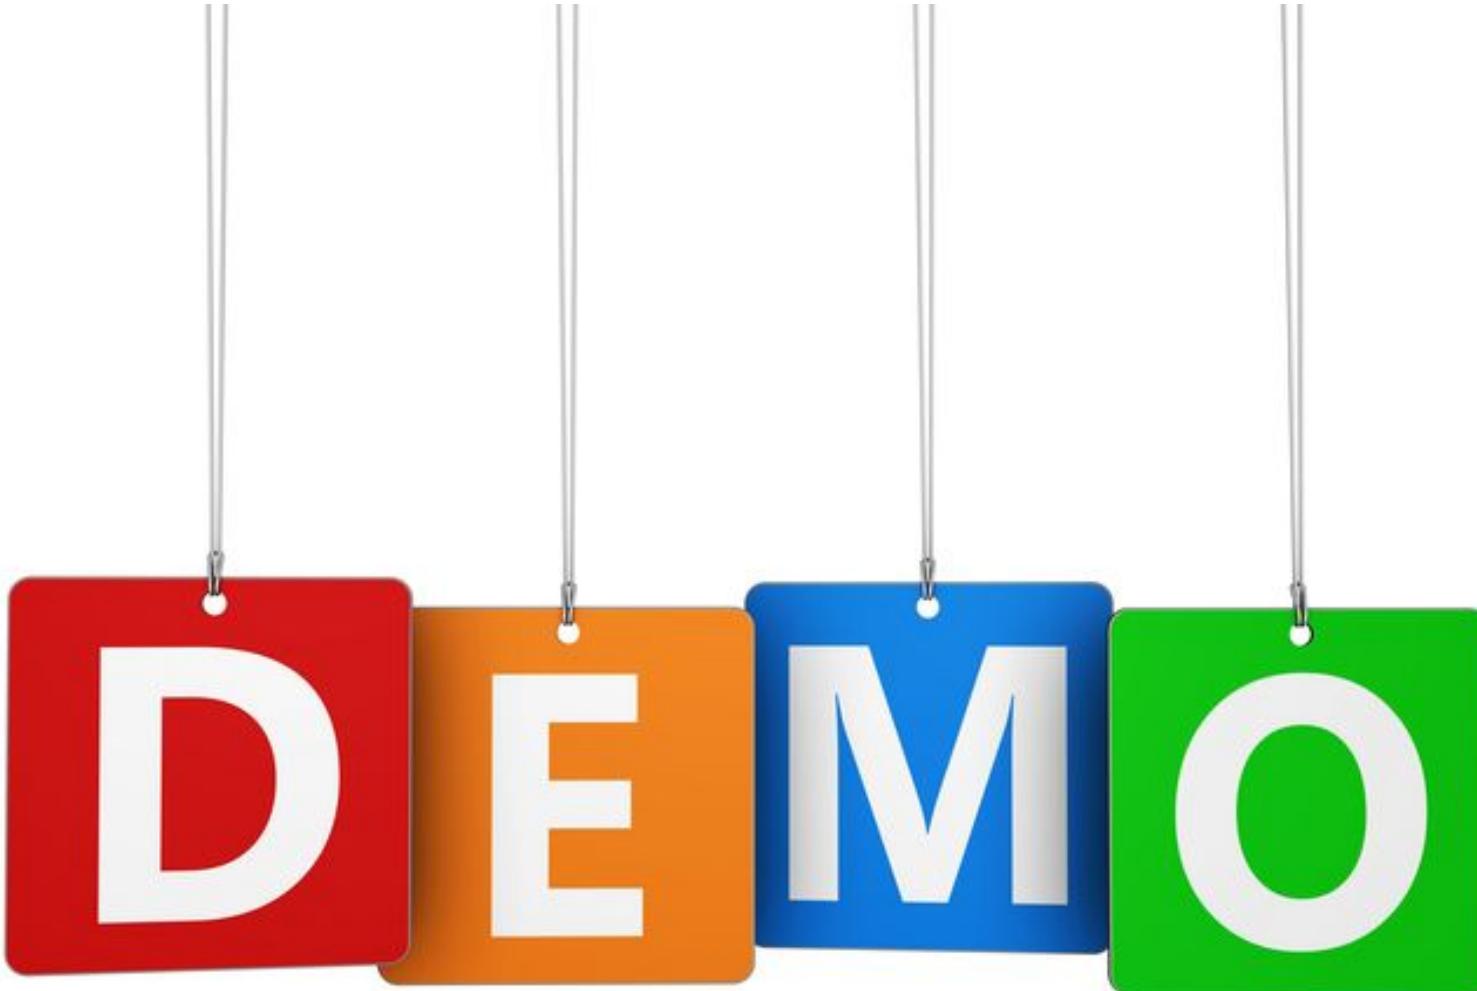


Containerization





Containerization

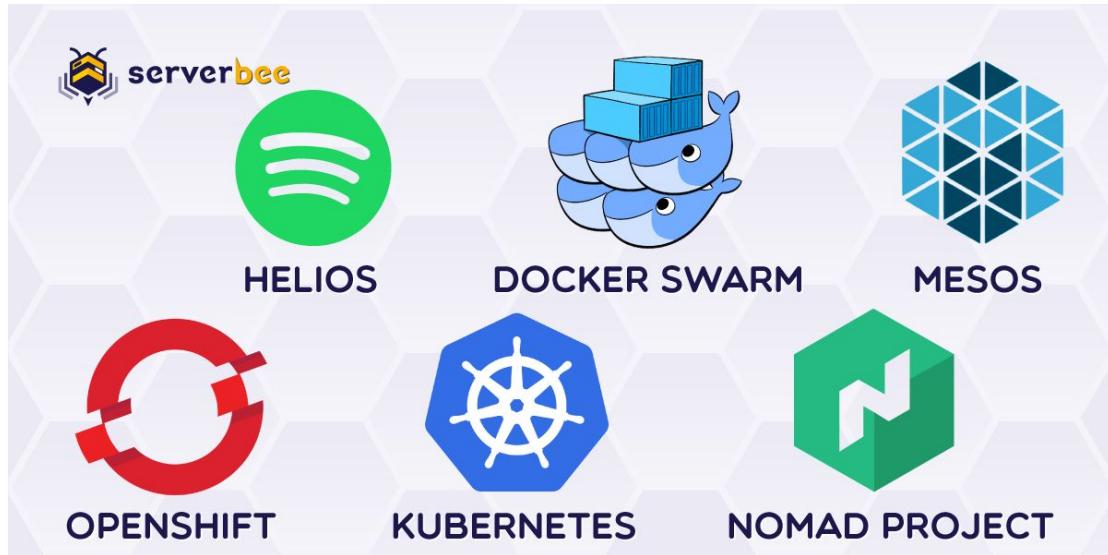


KUBERNETES

Kubernetes

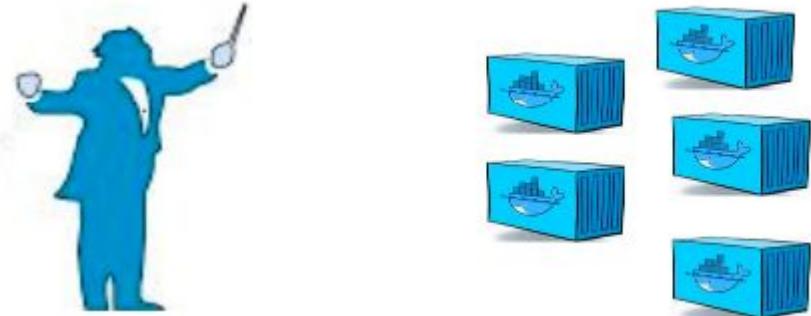
Orchestration

- Containers are great, but when you get lots of them running, at some point, you need them all working together in harmony to solve business problems.
- Tools to manage, scale, and maintain containerized applications are called orchestrators, and the most common example of this is **Kubernetes**.



Kubernetes

Orchestration



Container orchestration is used to automate the following tasks at scale:

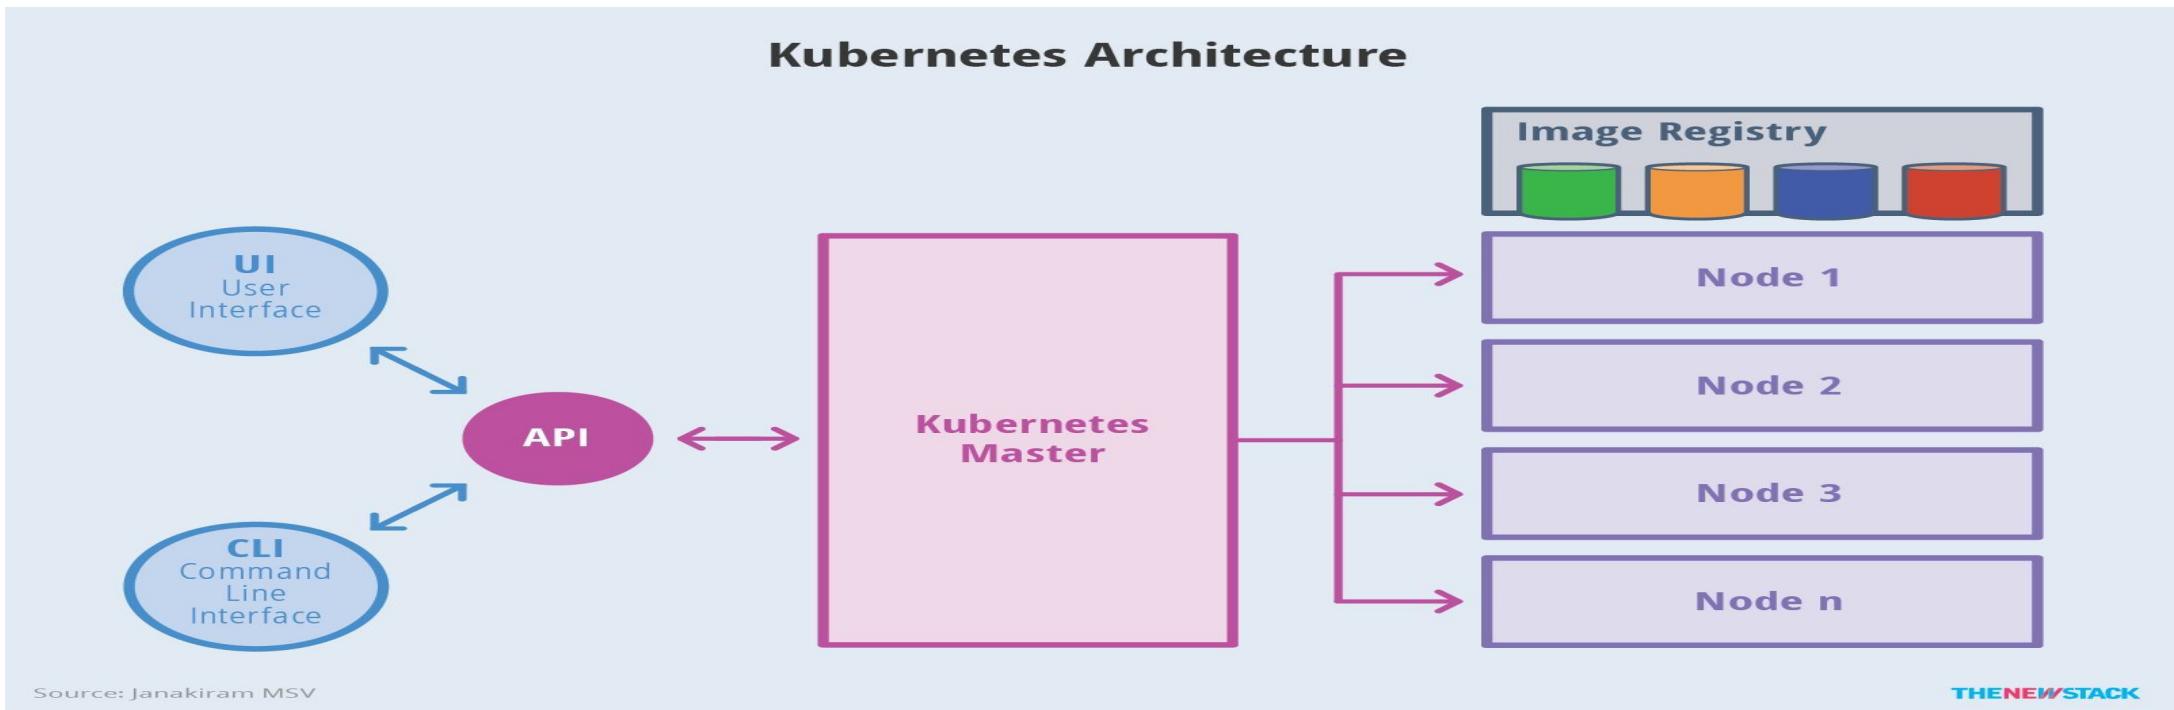
- Provisioning and deployments of containers
- Availability of containers
- Load balancing, traffic routing and service discovery of containers
- Health monitoring of containers
- Securing the interactions between containers.
- Configuring and scheduling of containers
- Allocation of resources between containers

Kubernetes

Kubernetes Components

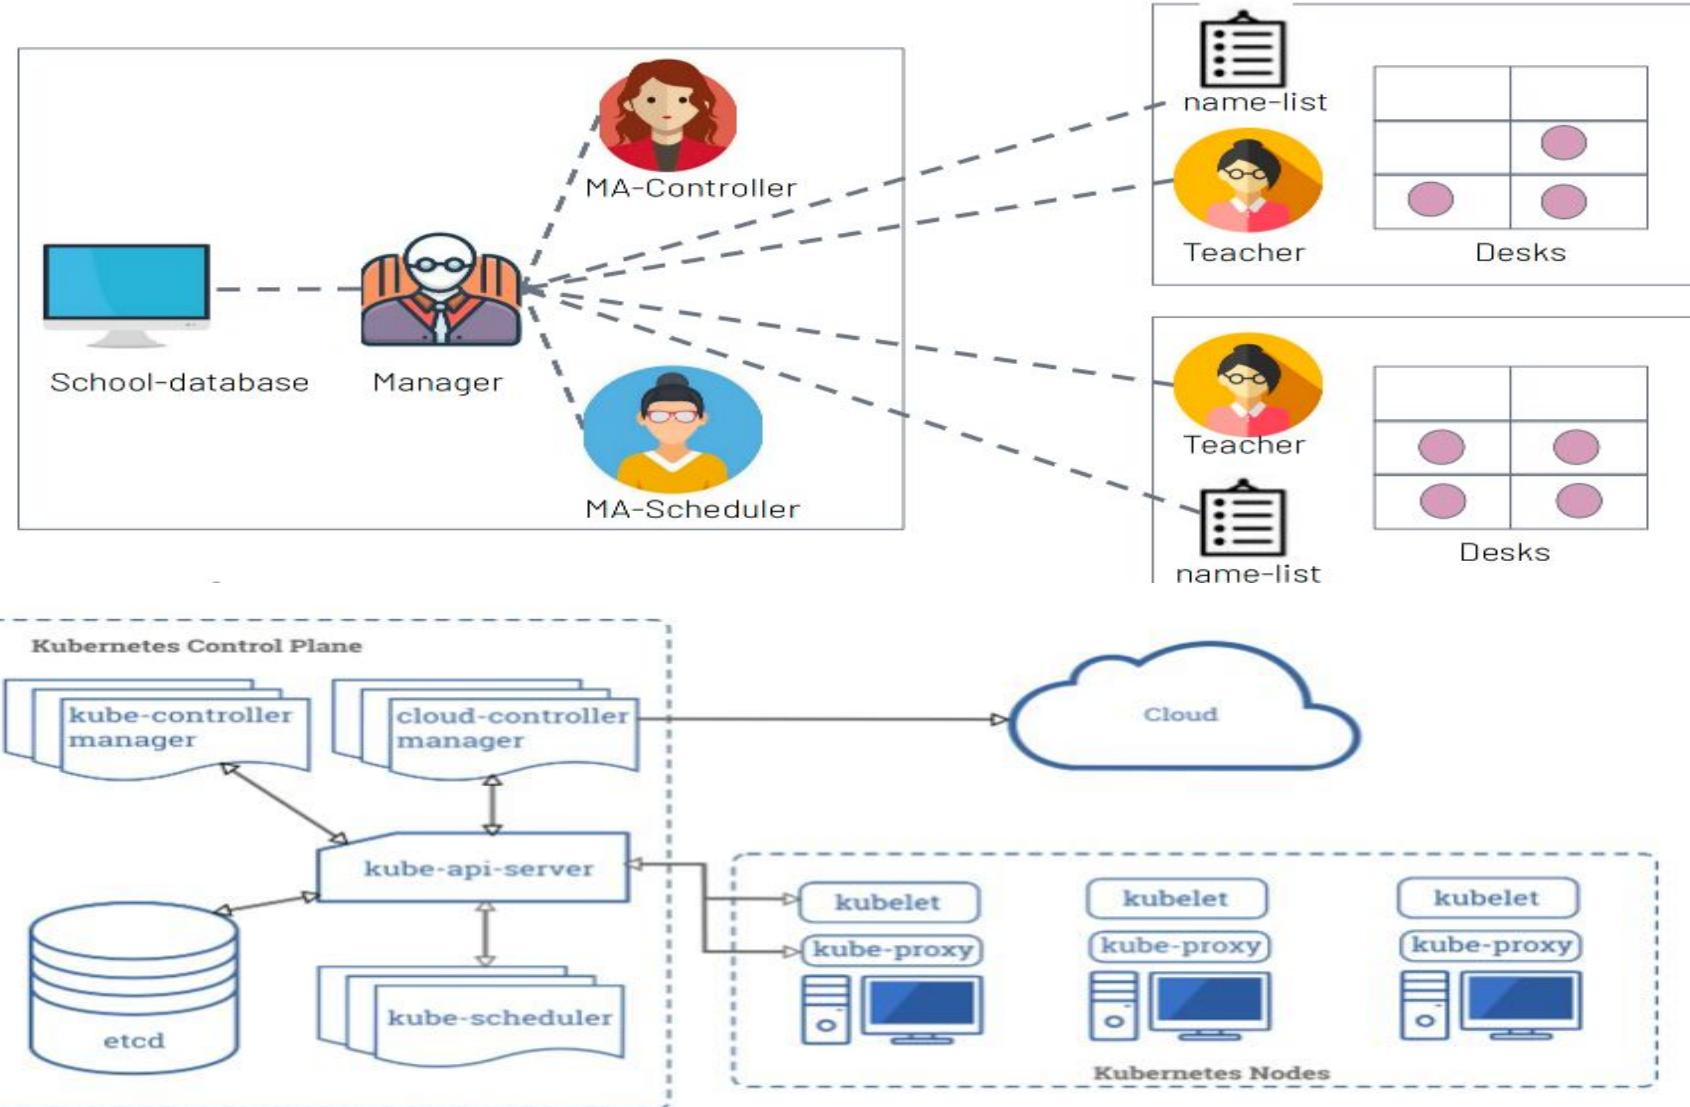
Kubernetes has the following main components:

- One or more master nodes
- One or more worker nodes.



Kubernetes

Kubernetes Components



Kubernetes

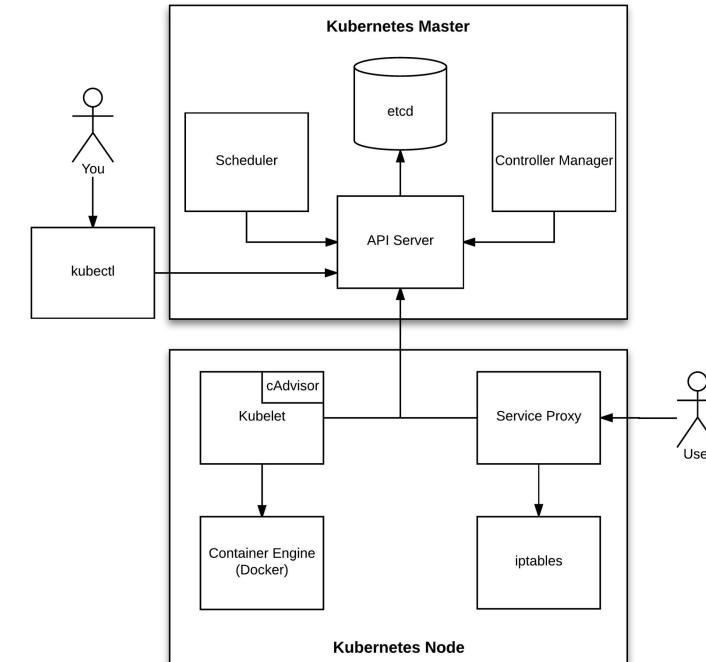
Kubernetes Components

kube-apiserver:

- Provides a forward facing REST interface into the kubernetes control plane and datastore.
- All clients and other applications interact with kubernetes strictly through the API Server.
- Acts as the gatekeeper to the cluster by handling authentication and authorization, request validation, mutation, and admission control in addition to being the front-end to the backing datastore.

etcd:

- etcd acts as the cluster datastore.
- Purpose in relation to Kubernetes is to provide a strong, consistent and highly available key-value store for persisting cluster state.
- Stores objects and config information.



Kubernetes

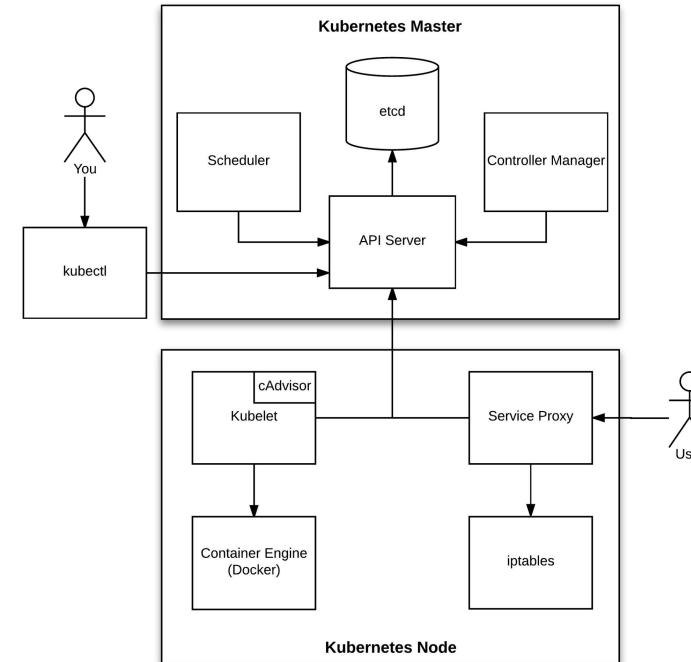
Kubernetes Components

kube-controller-manager:

- Serves as the primary daemon that manages all core component control loops.
- Monitors the cluster state via the apiserver and steers the cluster towards the desired state

kube-scheduler:

- Verbose policy-rich engine that evaluates workload requirements and attempts to place it on a matching resource.
- Workload Requirements can include: general hardware requirements, affinity/anti-affinity, labels, and other various custom resource requirements.



Kubernetes

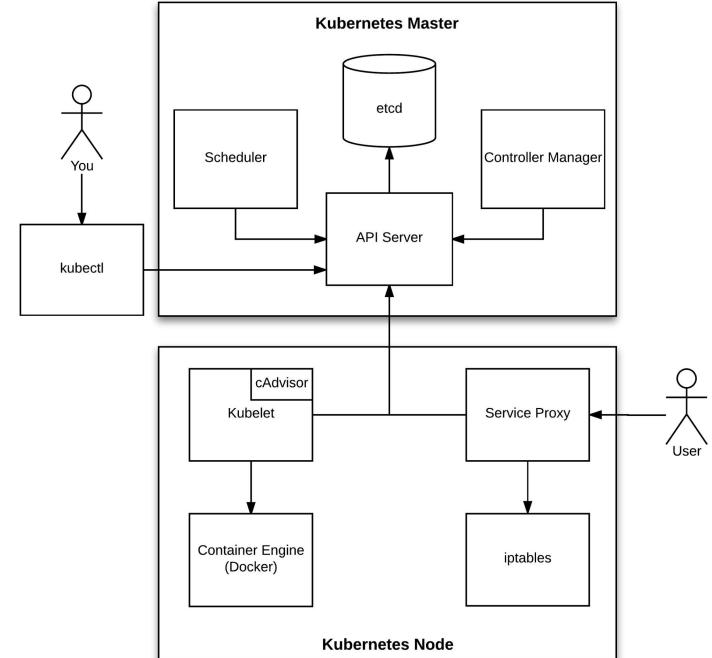
Kubernetes Components

kubelet:

- Acts as the node agent responsible for managing the lifecycle of every pod on its host.
- Kubelet understands YAML container manifests that it can read from several sources.

kube-proxy:

- Manages the network rules on each node.
- Performs connection forwarding or load balancing for Kubernetes cluster services.



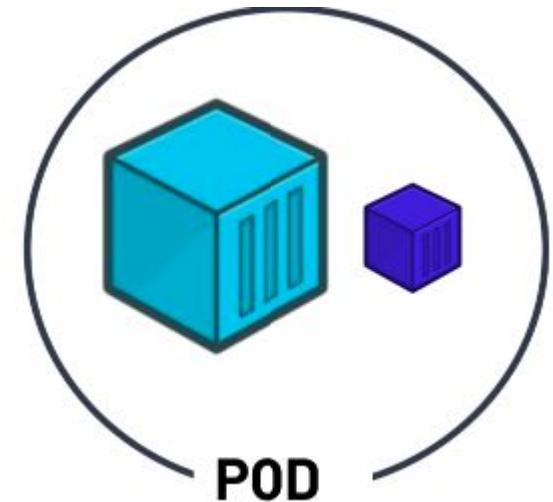
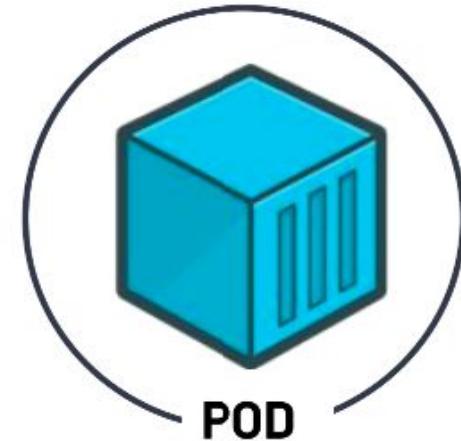
KUBERNETES Objects

- ▶ PODs
- ▶ ReplicaSets
- ▶ Deployment
- ▶ Namespaces
- ▶ Object Model

Kubernetes Objects

Pods

- Kubernetes doesn't deal with containers directly.
- PODs are Kubernetes objects that encapsulate the containers.
- Pods are the smallest deployable units of computing that you can create and manage in Kubernetes.
- A POD can have multiple containers.
- Sometimes an application need a helper container, such as logging, monitoring, etc.
- These helper containers should coexist with your application container.



Kubernetes Objects

Pods

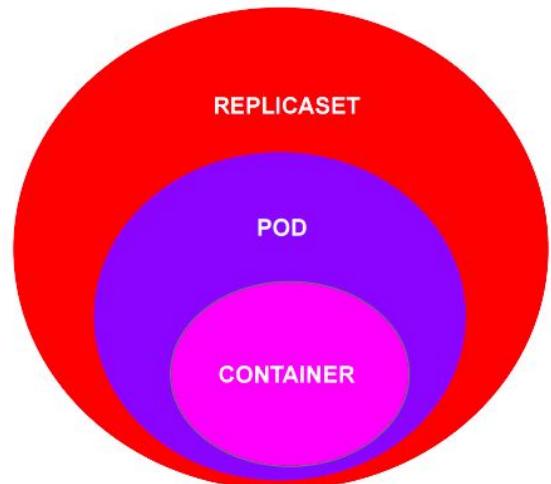
```
apiVersion: v1
kind: Pod
metadata:
  name: init-demo
spec:
  containers:
    - name: main-app
      image: nginx
  initContainers:
    - name: init-myservice
      image: busybox
      command: ['sh', '-c', 'echo Initializing Service; sleep 30']
```

```
apiVersion: v1
kind: Pod
metadata:
  name: sidecar-demo
spec:
  containers:
    - name: main-app
      image: nginx
    - name: logger-sidecar
      image: fluentd
      volumeMounts:
        - name: log-volume
          mountPath: /var/log/nginx
  volumes:
    - name: log-volume
      emptyDir: {}
```

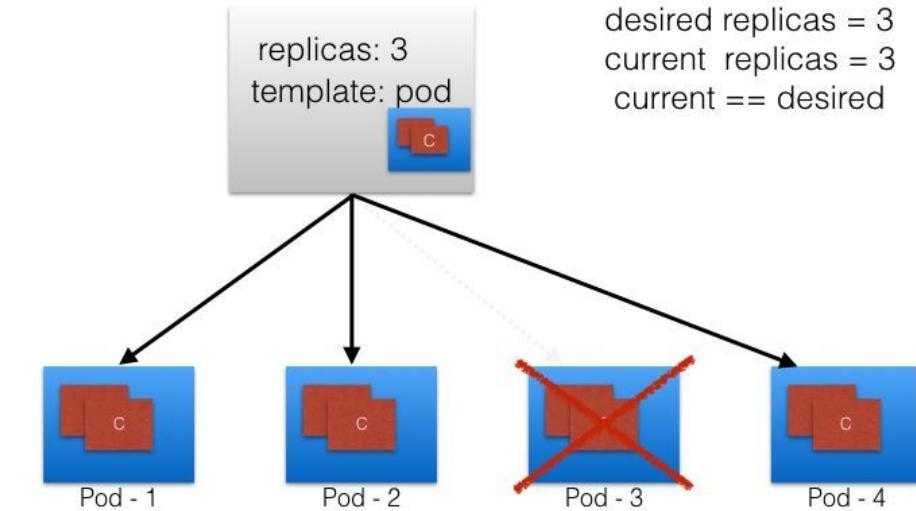
Kubernetes Objects

ReplicaSets

- A **ReplicaSet's** purpose is to maintain a stable set of replica Pods running at any given time.
- Even if you have a single POD, the ReplicaSet will bring up a new POD when the existing one fails.



Replica Set



Kubernetes Objects

ReplicaSets

Pod to ReplicaSet

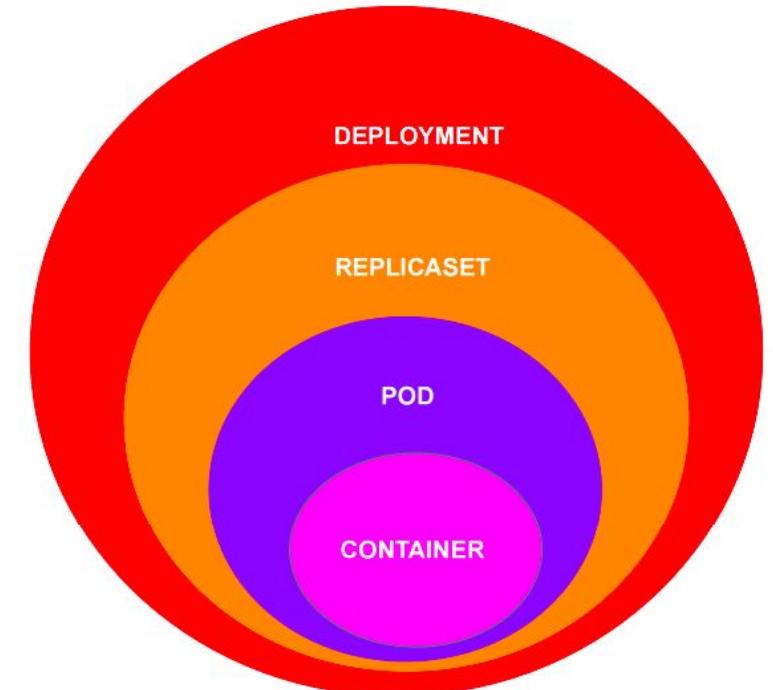
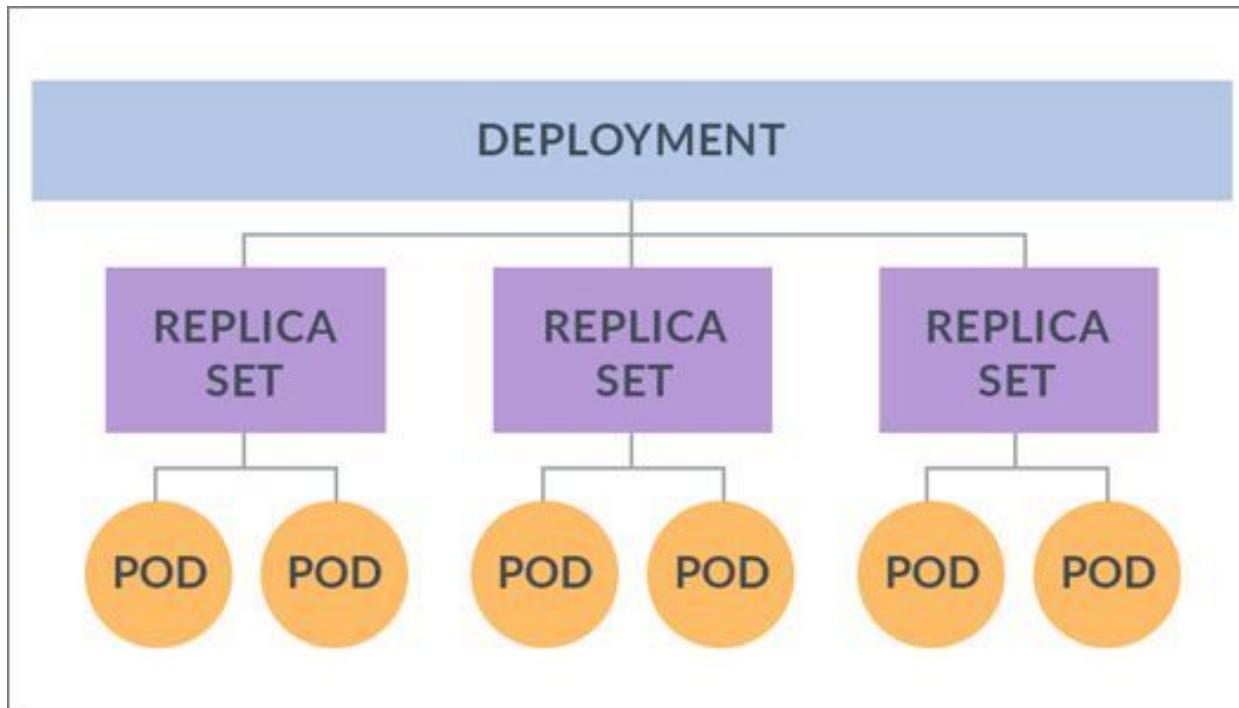
```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: nginx
spec:
  containers:
  - name: mynginx
    image: nginx:1.19
    ports:
    - containerPort: 80
```

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-rs
  labels:
    environment: dev
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: mynginx
        image: nginx:1.19
        ports:
        - containerPort: 80
```

Kubernetes Objects

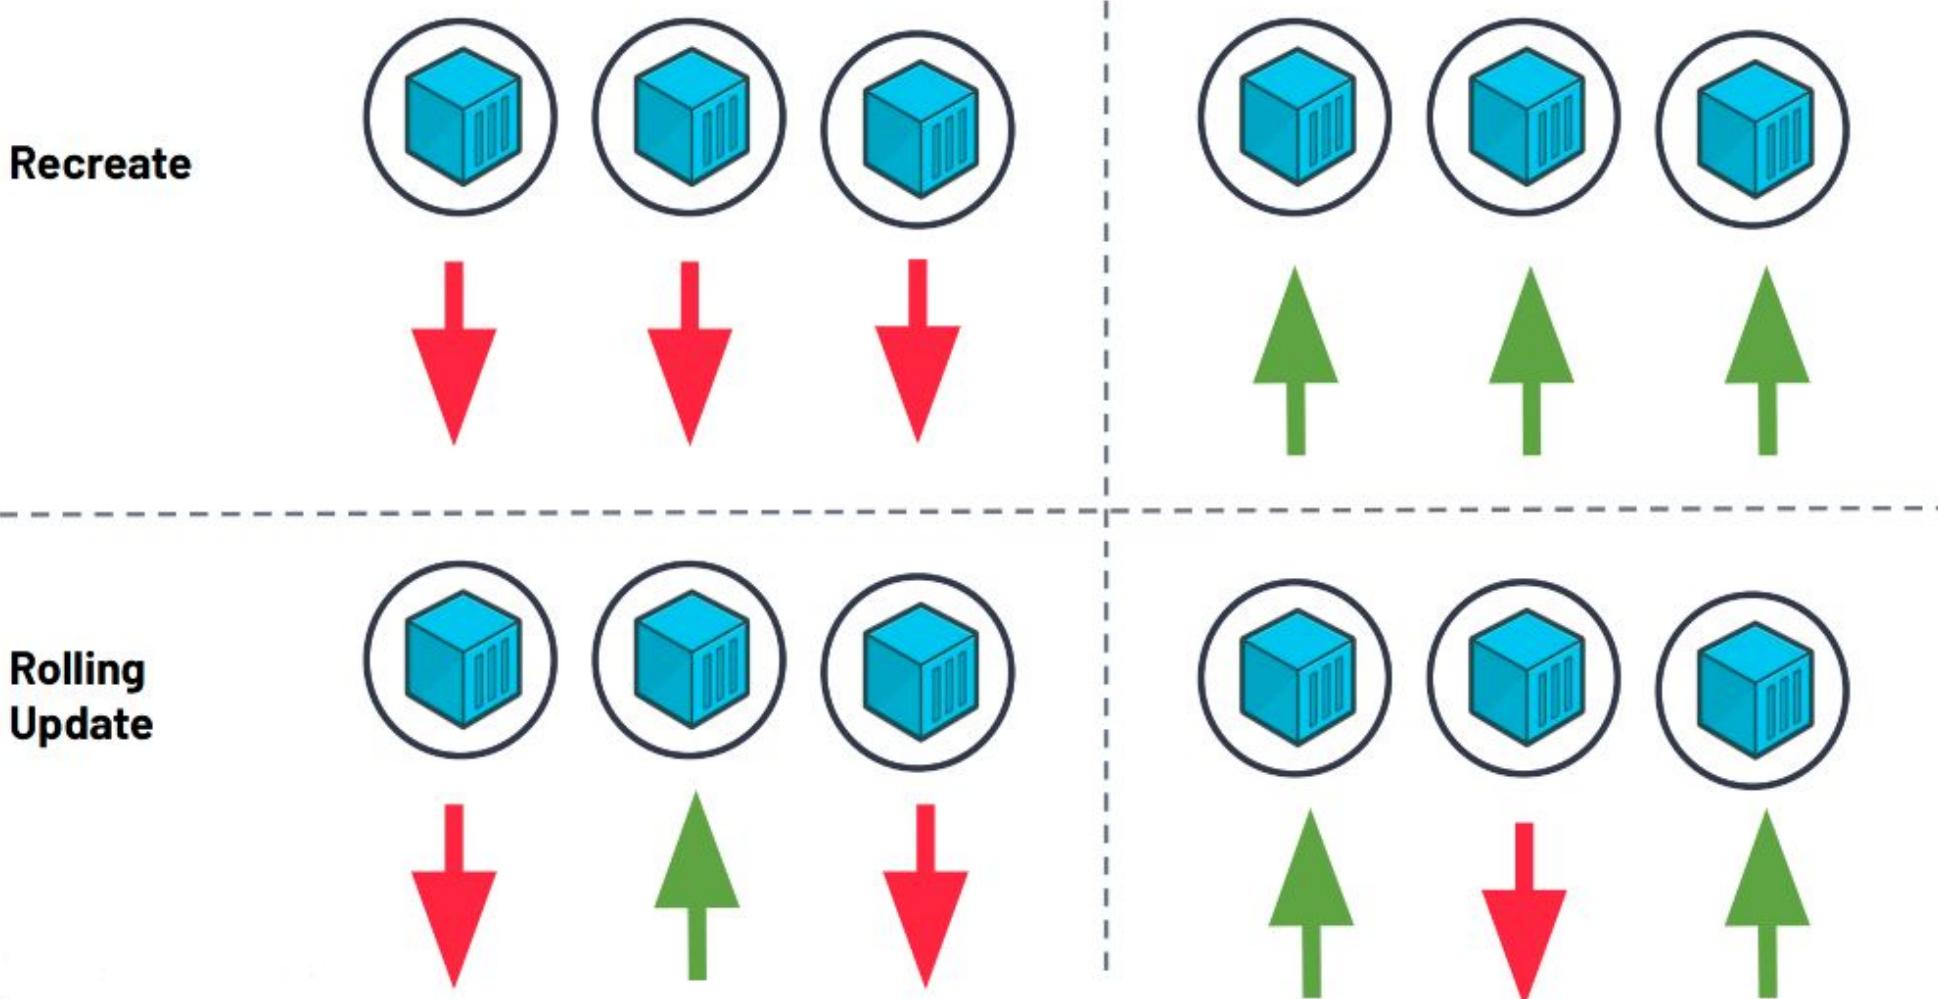
Deployment

- One step higher in the hierarchy, deployments provides declarative updates for Pods and ReplicaSets.



Kubernetes Objects

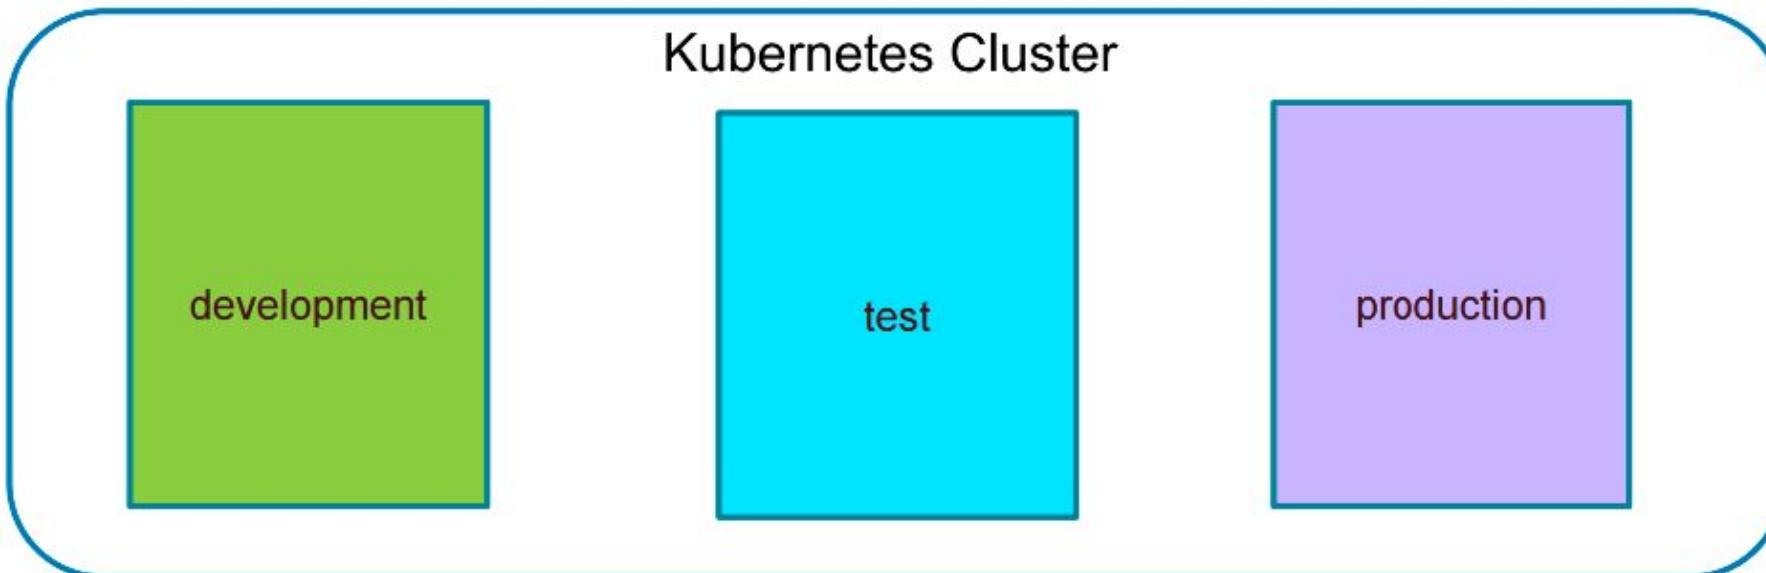
Deployment Strategy



Kubernetes Objects

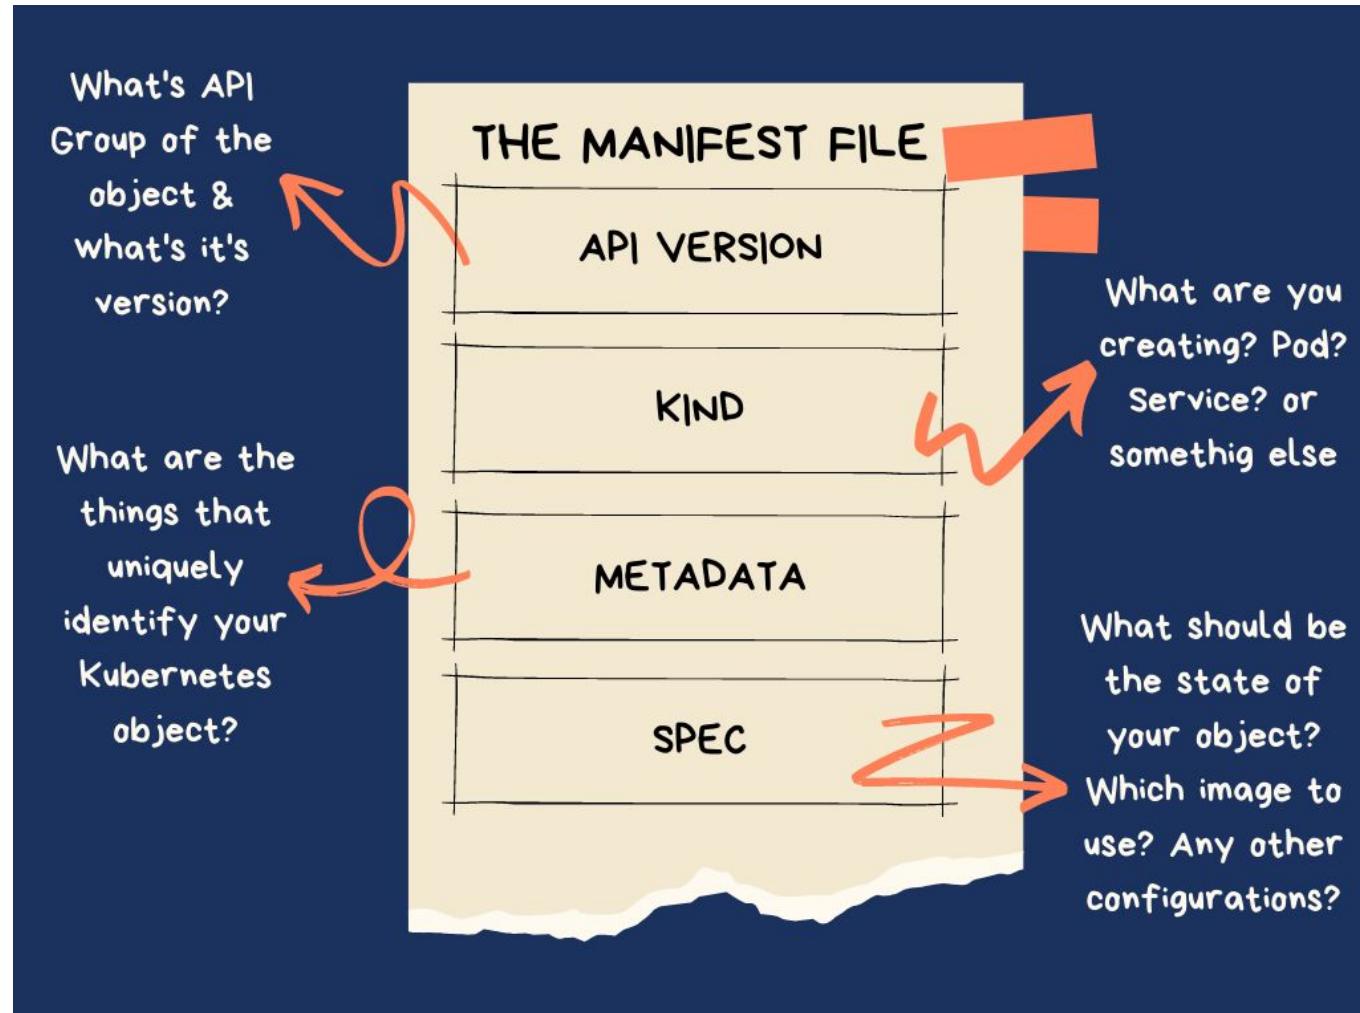
Namespaces

- Kubernetes supports multiple virtual clusters backed by the same physical cluster. These virtual clusters are called **namespaces**.
- Namespaces are intended for use in environments with many users spread across multiple teams, or projects.



Kubernetes Objects

Object Model



Kubernetes Objects

Object Model

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
        ports:
          - containerPort: 80
```

- All objects must have **apiVersion**, **kind**, **metadata** and **spec** fields.
- **apiVersion:** Which version of the Kubernetes API you're using to create this object
- **kind:** What kind of object you want to create
- **metadata:** Data that helps uniquely identify the object, including a **name** string, **labels**, and optional **namespace**
- **spec:** What state you desire for the object

KUBERNETES Networking

- ▶ Cluster Networking
- ▶ Services
- ▶ Service Types
- ▶ Ingress

Kubernetes Networking

Cluster Networking

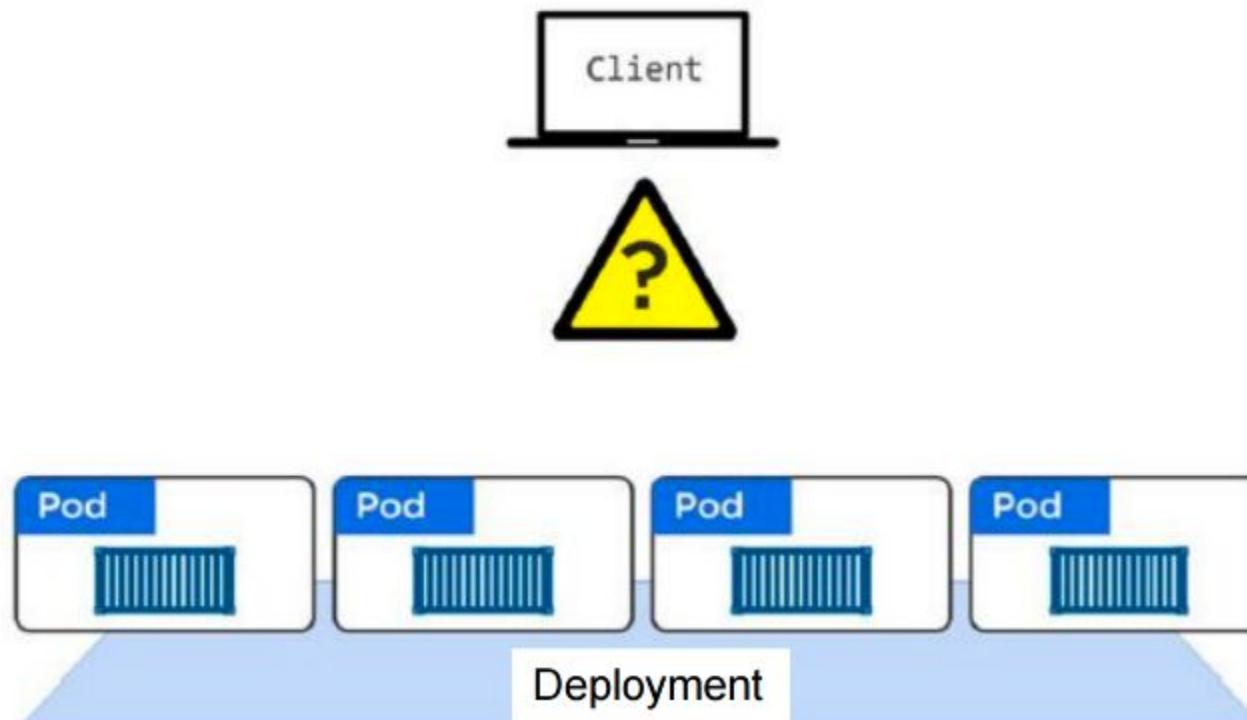
There are 4 distinct networking problems to address:

1. container-to-container communications:
2. Pod-to-Pod communications:
3. Pod-to-Service communications: this is covered by services.
4. External-to-Service communications: this is covered by services.

Kubernetes Networking

Services

Pods are not reliable

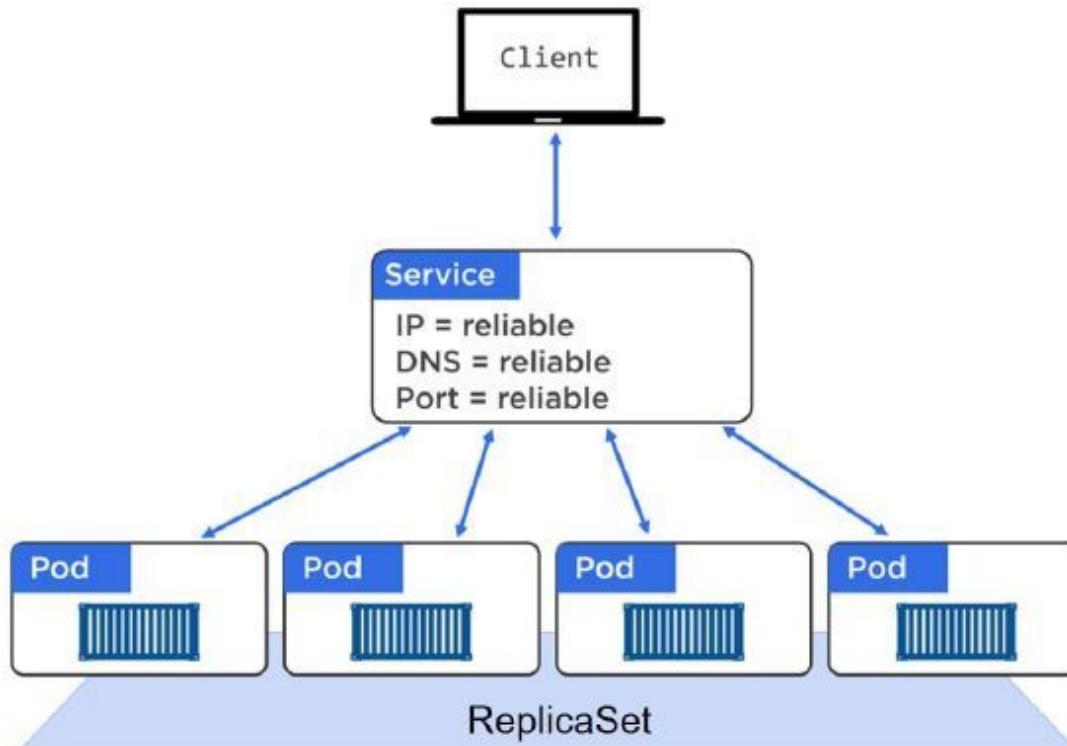


Kubernetes Networking

Services

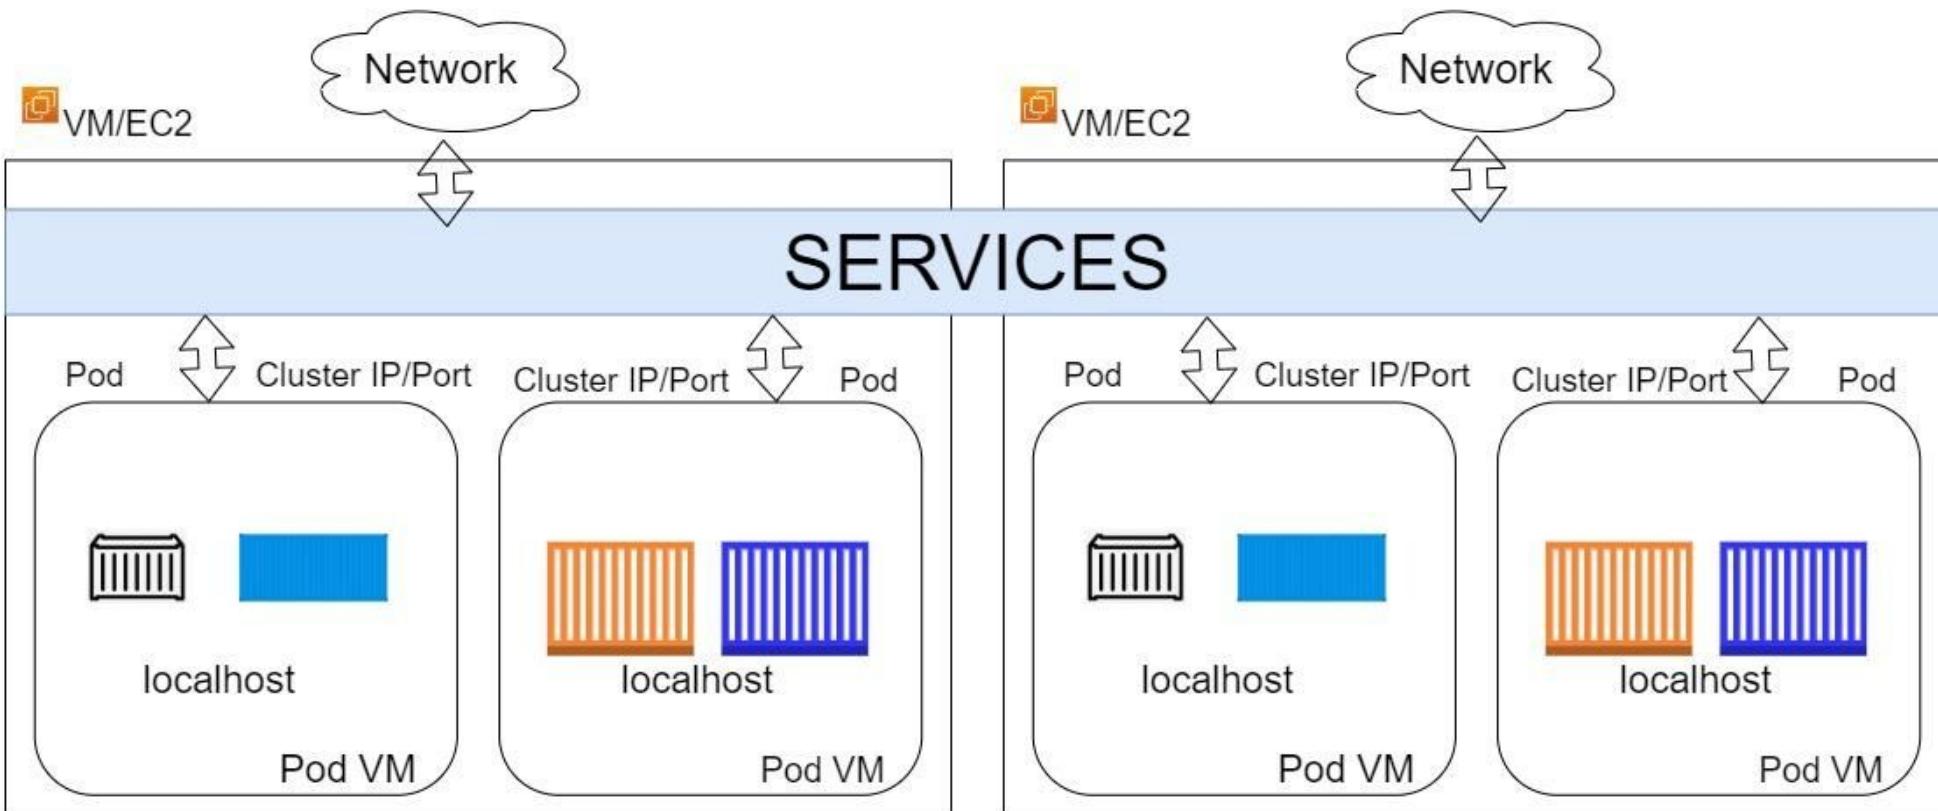
A **Service** offers a single **DNS entry** for a containerized application managed by the Kubernetes cluster

The Service is associated with the Pods, and provides them with a stable IP, DNS and port. It also loadbalances requests across the Pods.



Kubernetes Networking

Services

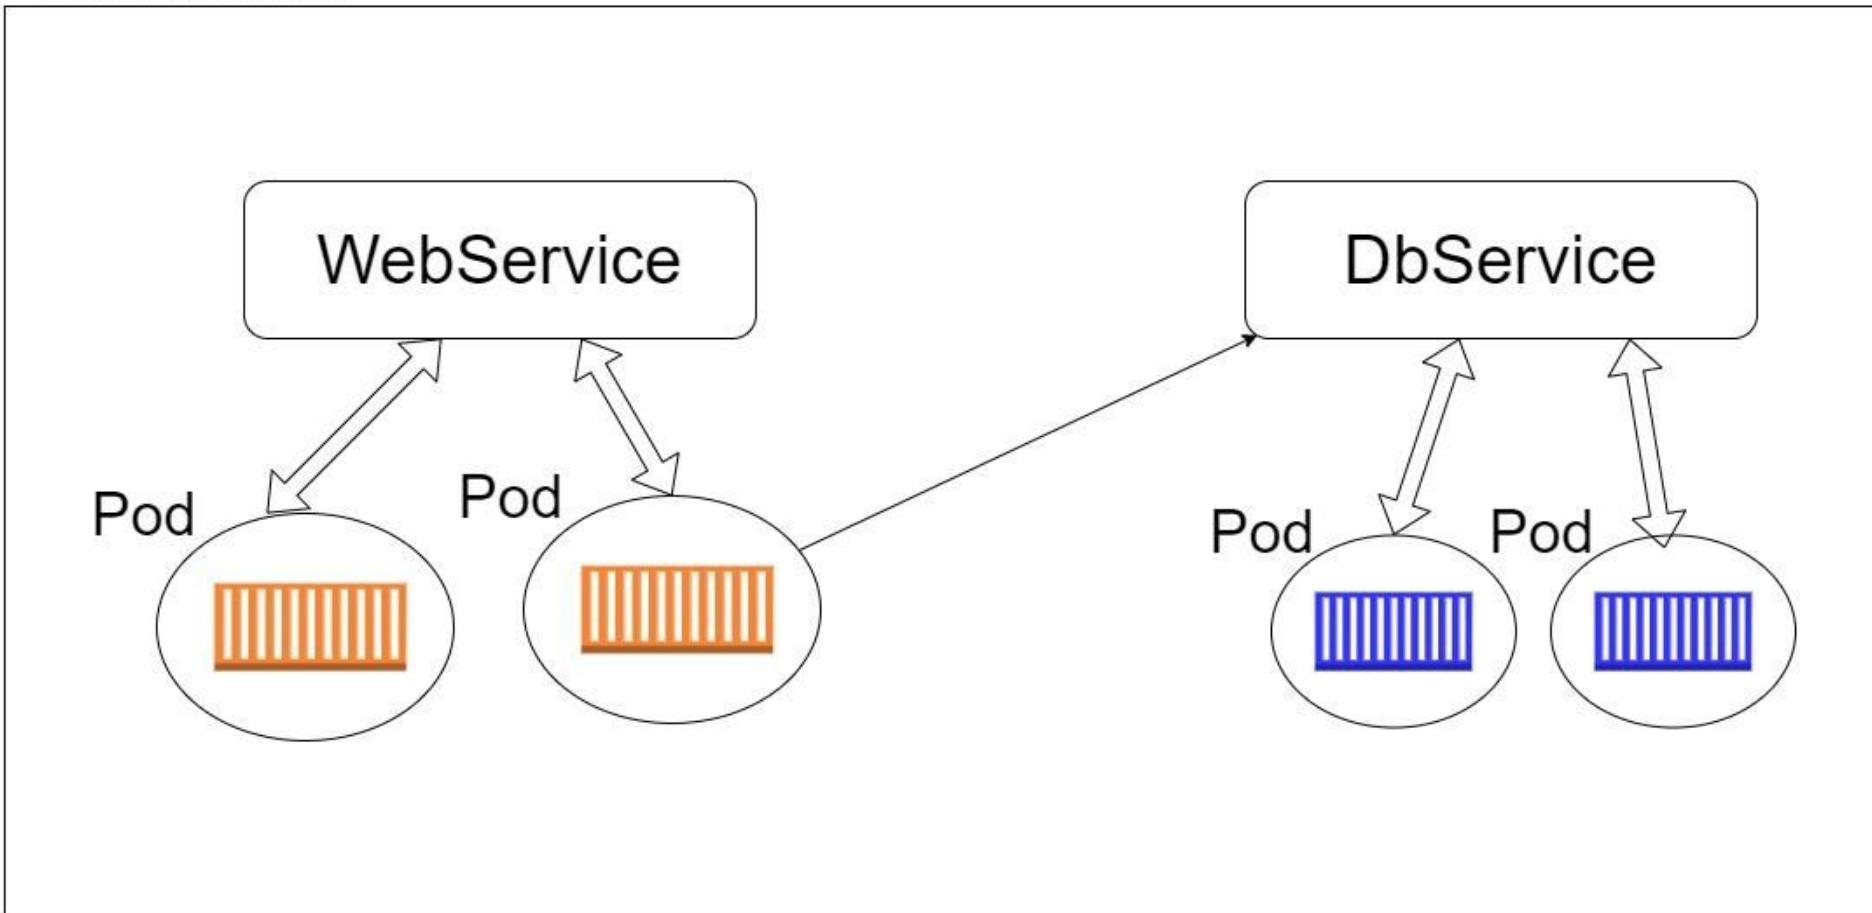


K8s Network

Kubernetes Networking

Services

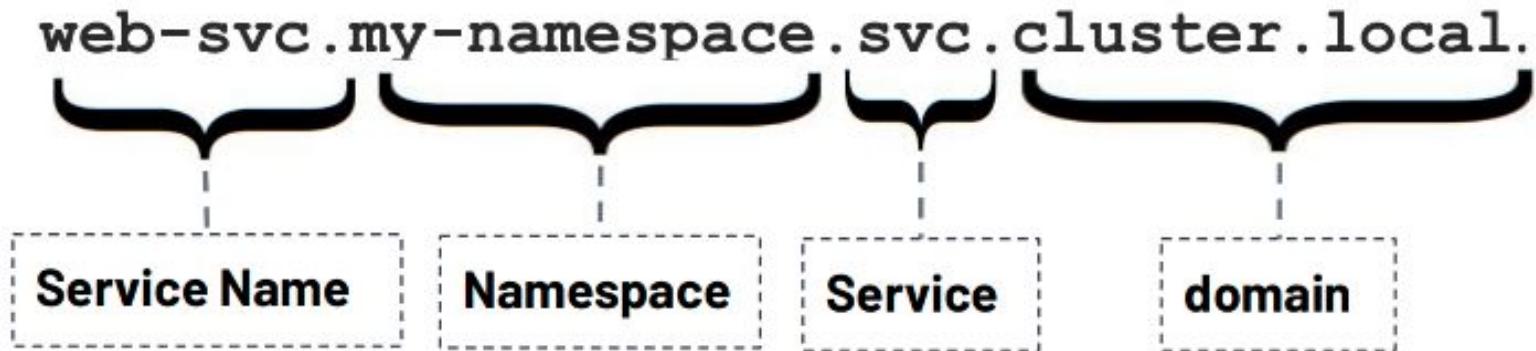
K8s Cluster



Kubernetes Networking

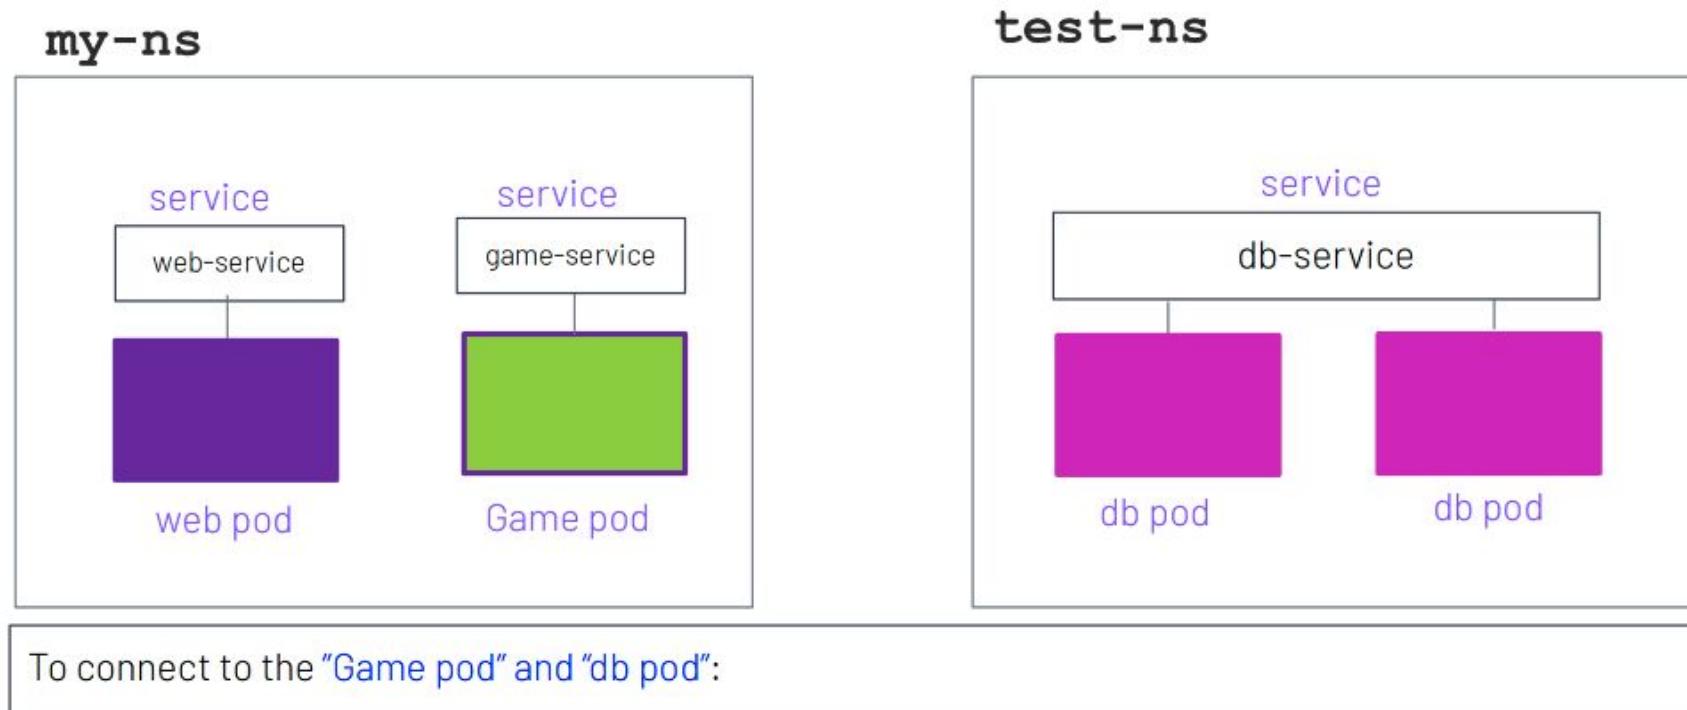
Services Discovery

Kubernetes has an add-on for **DNS**, which creates a DNS record for each Service and its format is



Kubernetes Networking

Services Discovery



From " web pod" -> " Game pod" --> hostname: game-service.my-ns:port
game-service:port

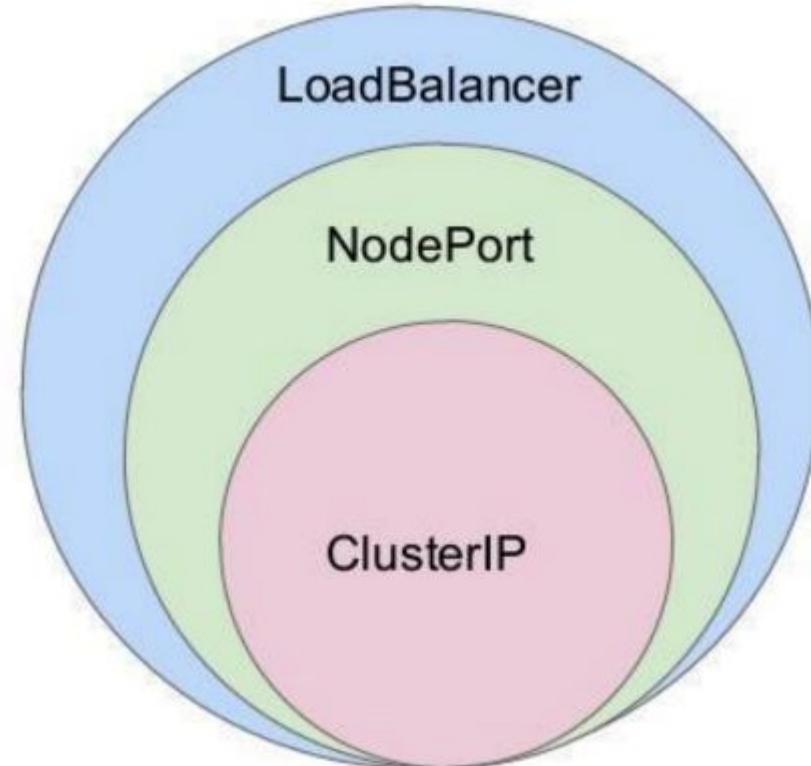
From " web pod" -> "db pod" --> hostname: db-service.test-ns.svc.cluster.local:port

Kubernetes Networking

Services Types

There are 3 major service types:

- ClusterIP (default)
- NodePort
- LoadBalancer



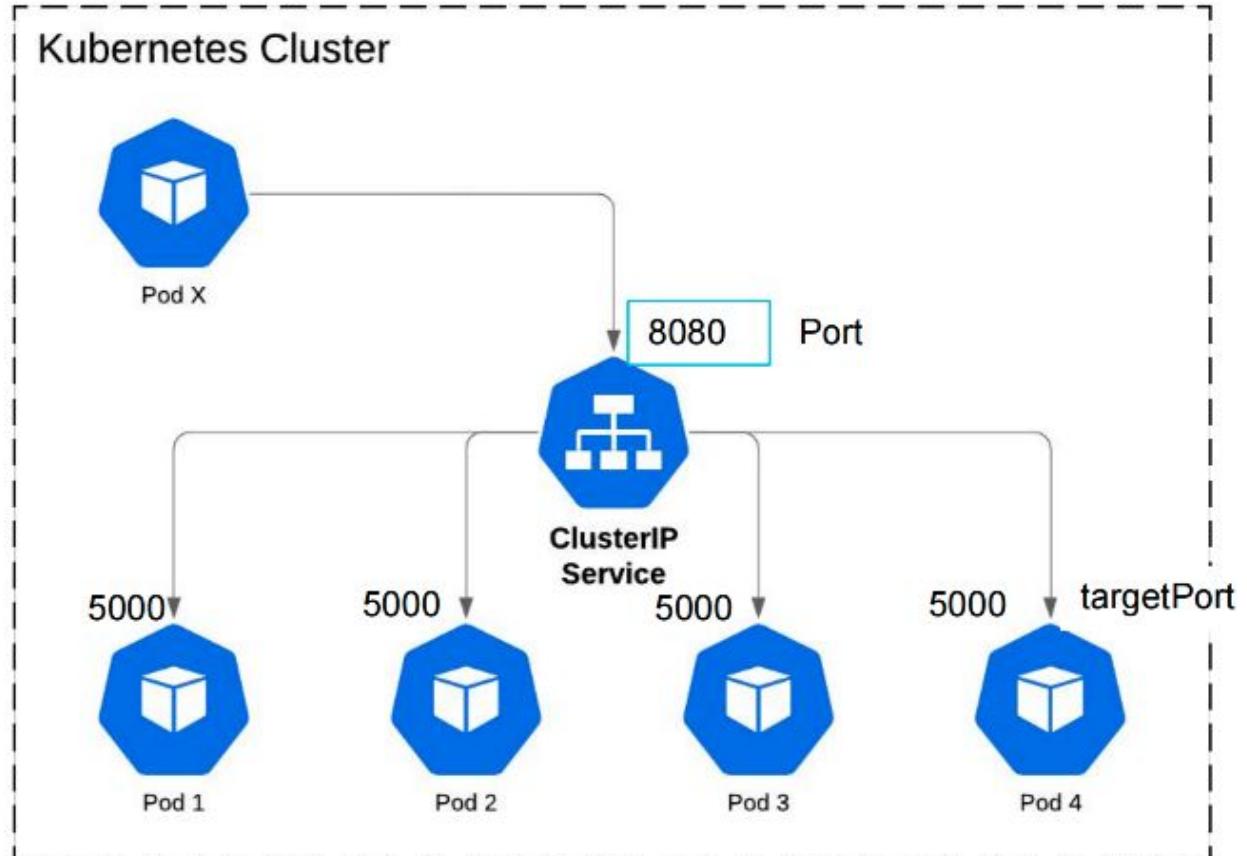
Kubernetes Networking

ClusterIP

Expose traffic internally

Example Usecase:

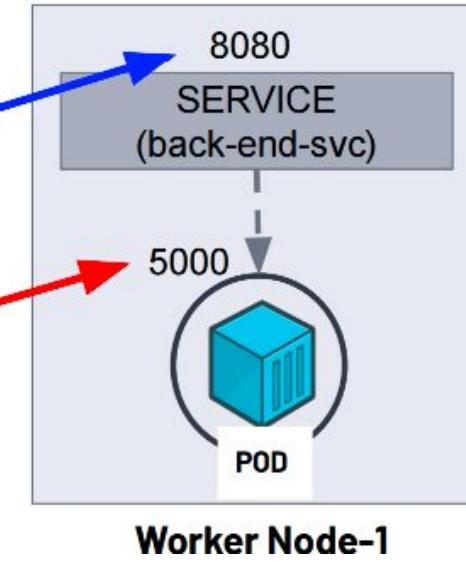
Good for service of frontend & back-end apps.



Kubernetes Networking

ClusterIP

```
apiVersion: v1
kind: Service
metadata:
  name: back-end-svc
  labels:
    app: back-end
spec:
  type: ClusterIP (default)
  selector:
    app: back-end
  ports:
  - port: 8080
    protocol: TCP
    targetPort: 5000
```



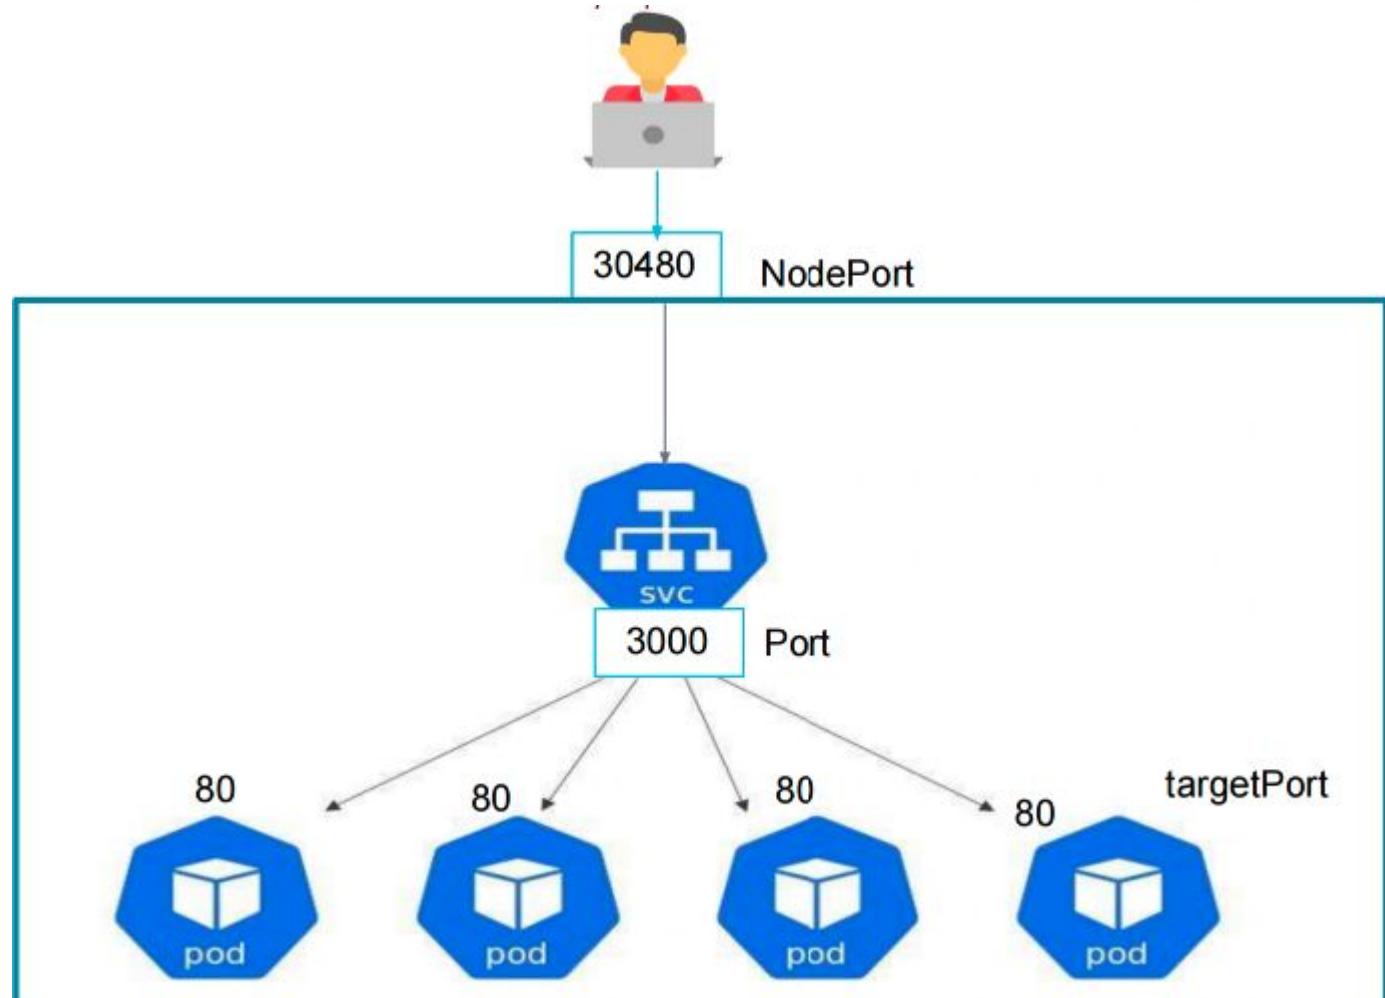
Kubernetes Networking

NodePort

Exposes traffic to the outside.

Example Use Case:

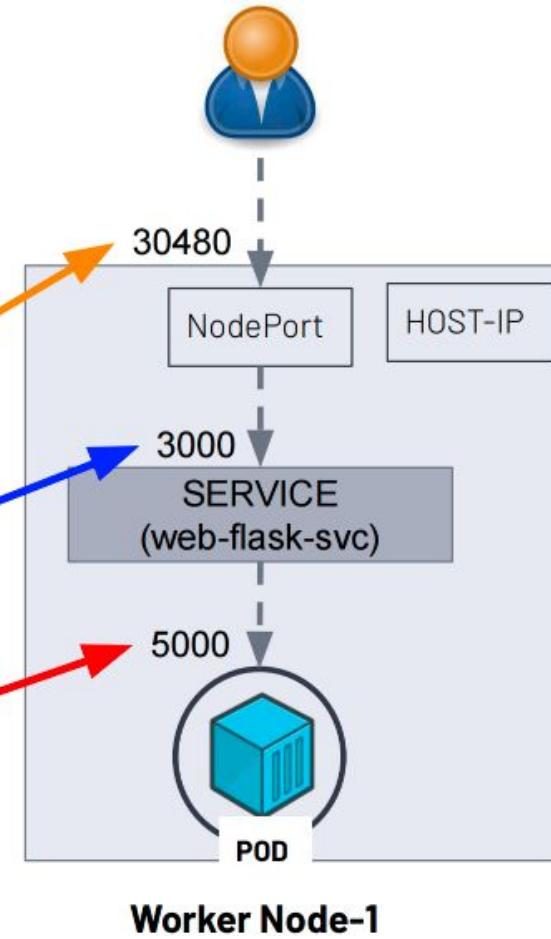
when we want to make our Services which has our app or website accessible from the external world.



Kubernetes Networking

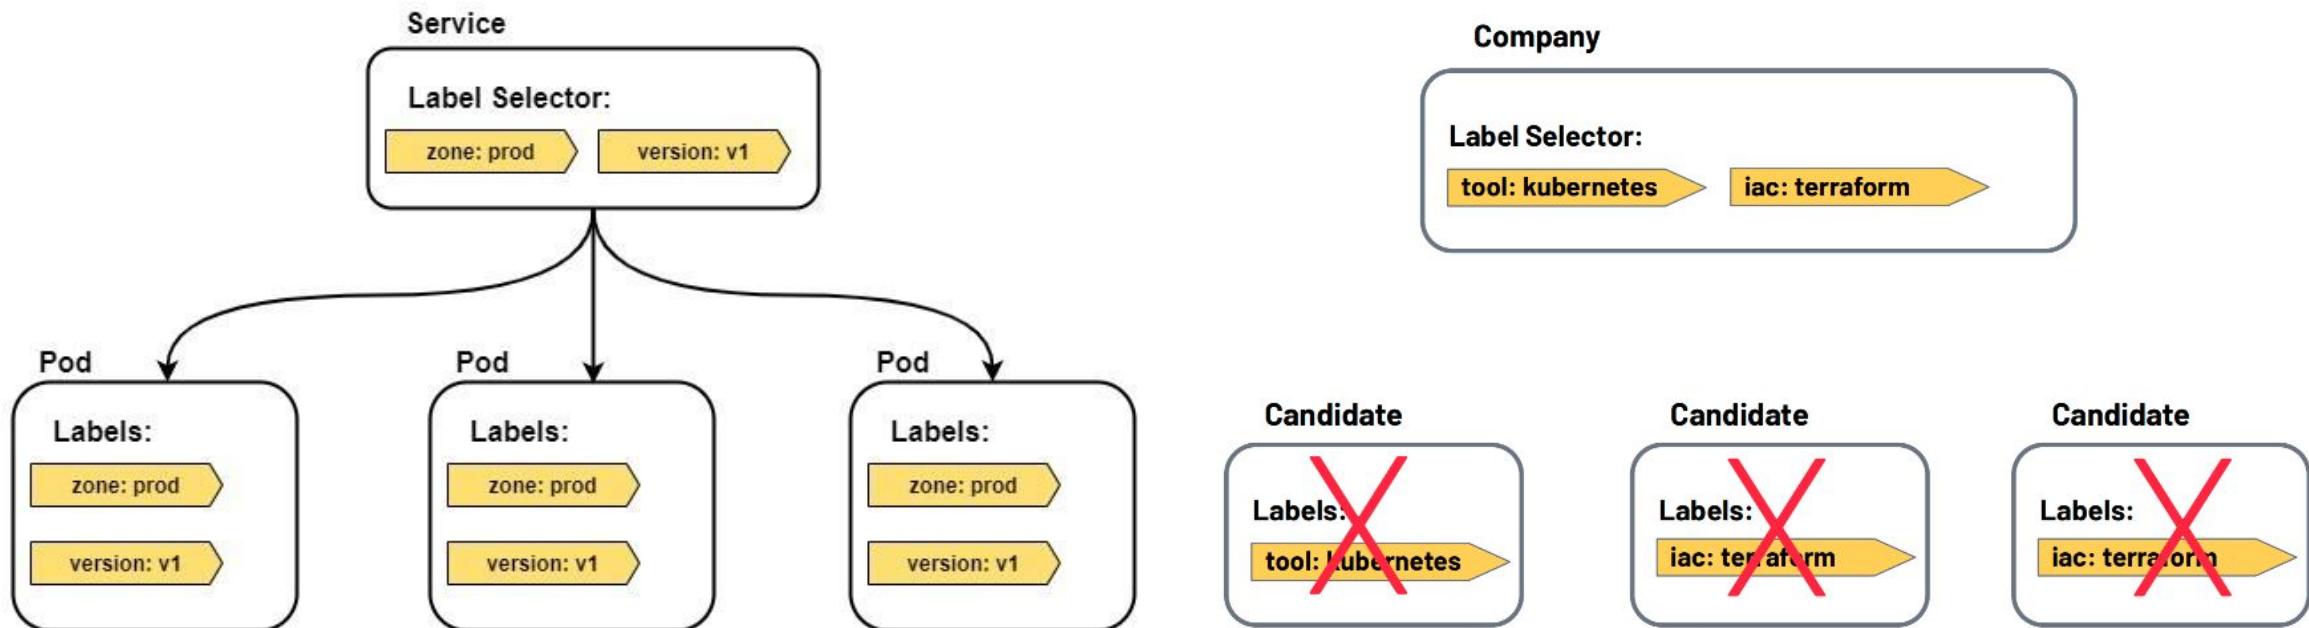
NodePort

```
apiVersion: v1
kind: Service
metadata:
  name: web-flask-svc
  labels:
    app: web-flask
spec:
  type: NodePort
  selector:
    app: web-flask
  ports:
    - nodePort: 30480
      port: 3000
      protocol: TCP
      targetPort: 5000
```



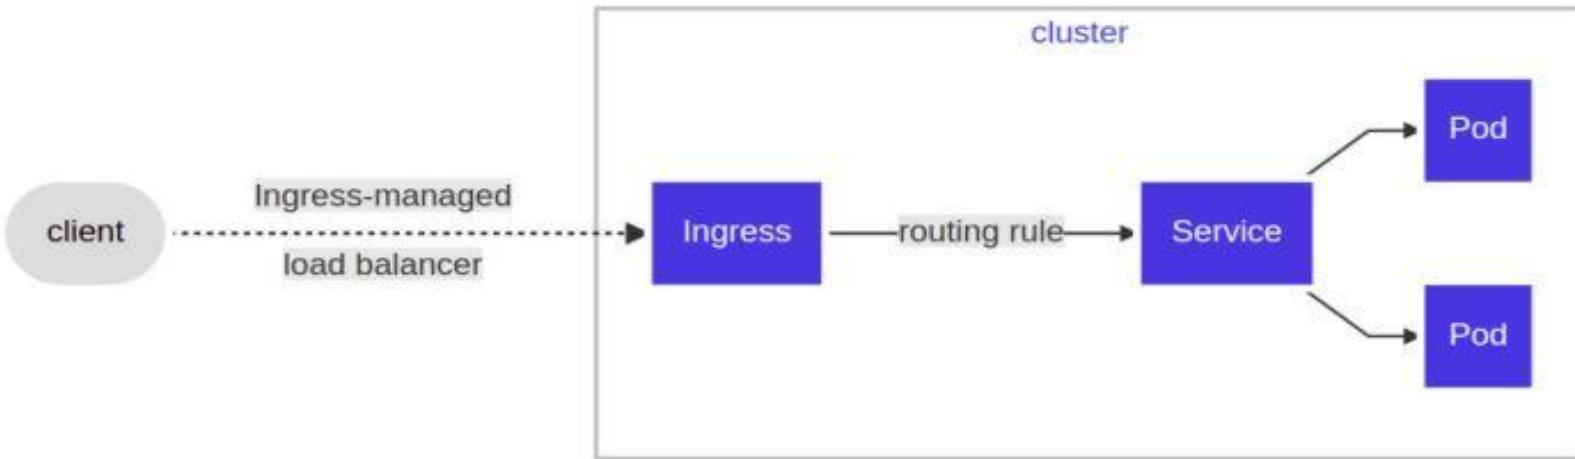
Kubernetes Networking

Labels and loose coupling



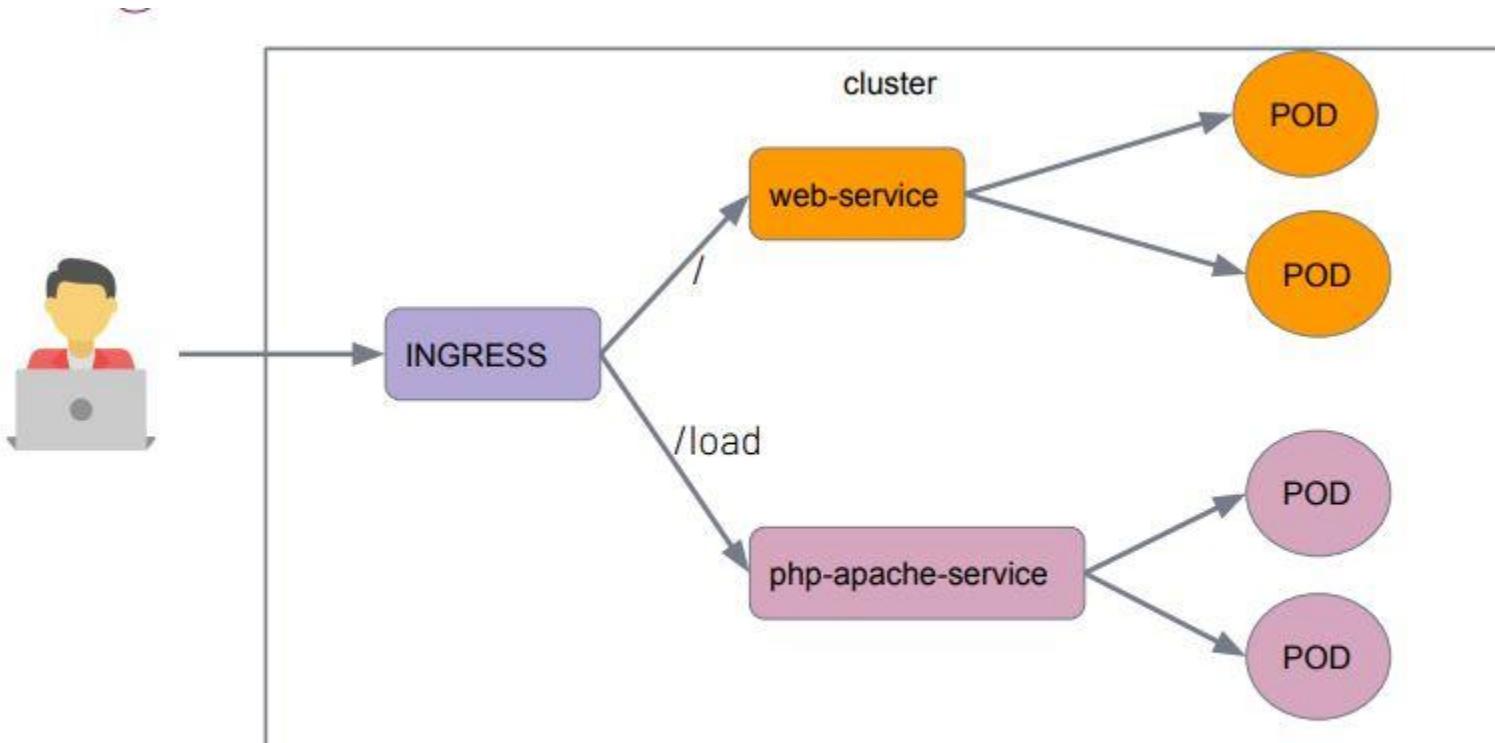
Kubernetes Networking

Ingress: With Ingress, users do not connect directly to a Service. Users reach the Ingress endpoint, and, from there, the request is forwarded to the desired Service.



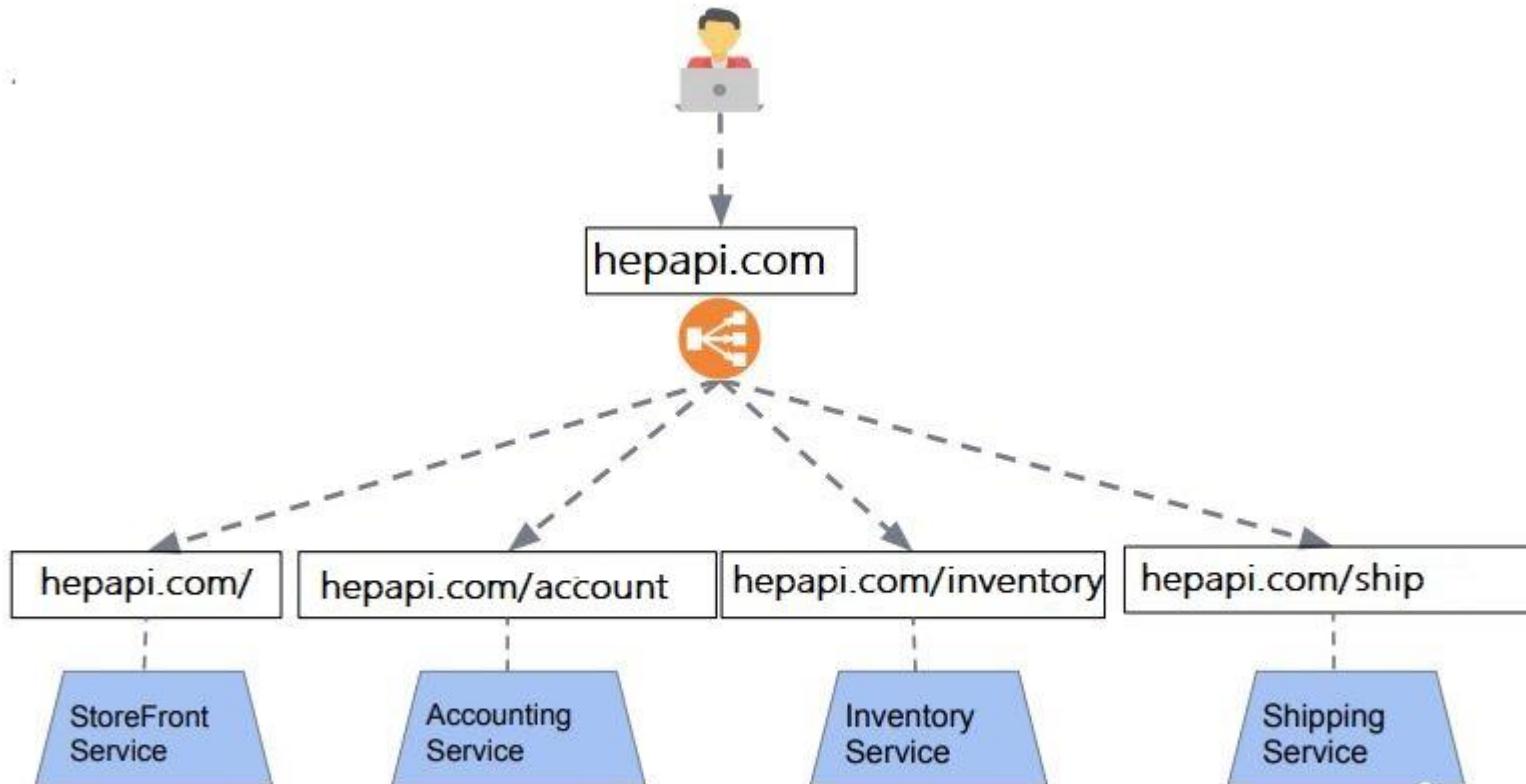
Kubernetes Networking

Ingress



Kubernetes Networking

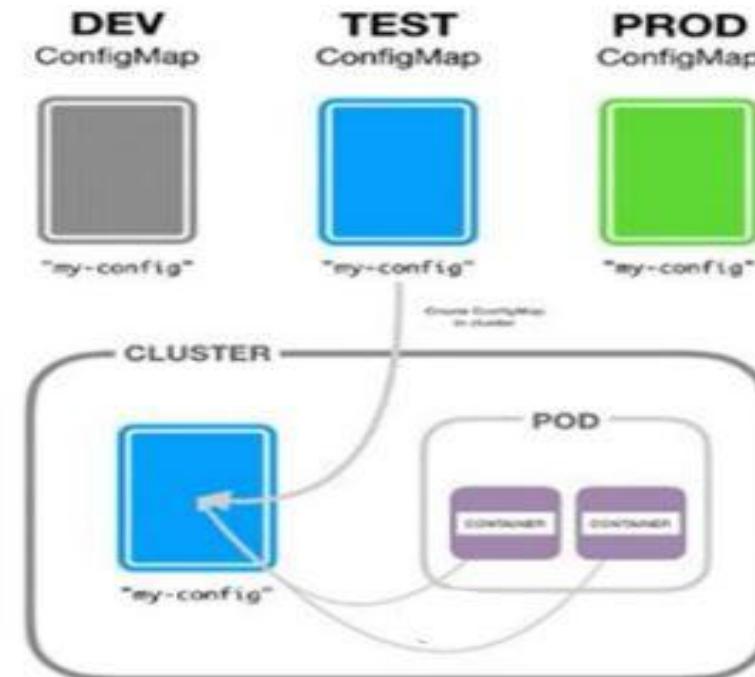
Ingress: exposes HTTP and HTTPS routes from outside the cluster to services within the cluster.



Secrets and Configmaps

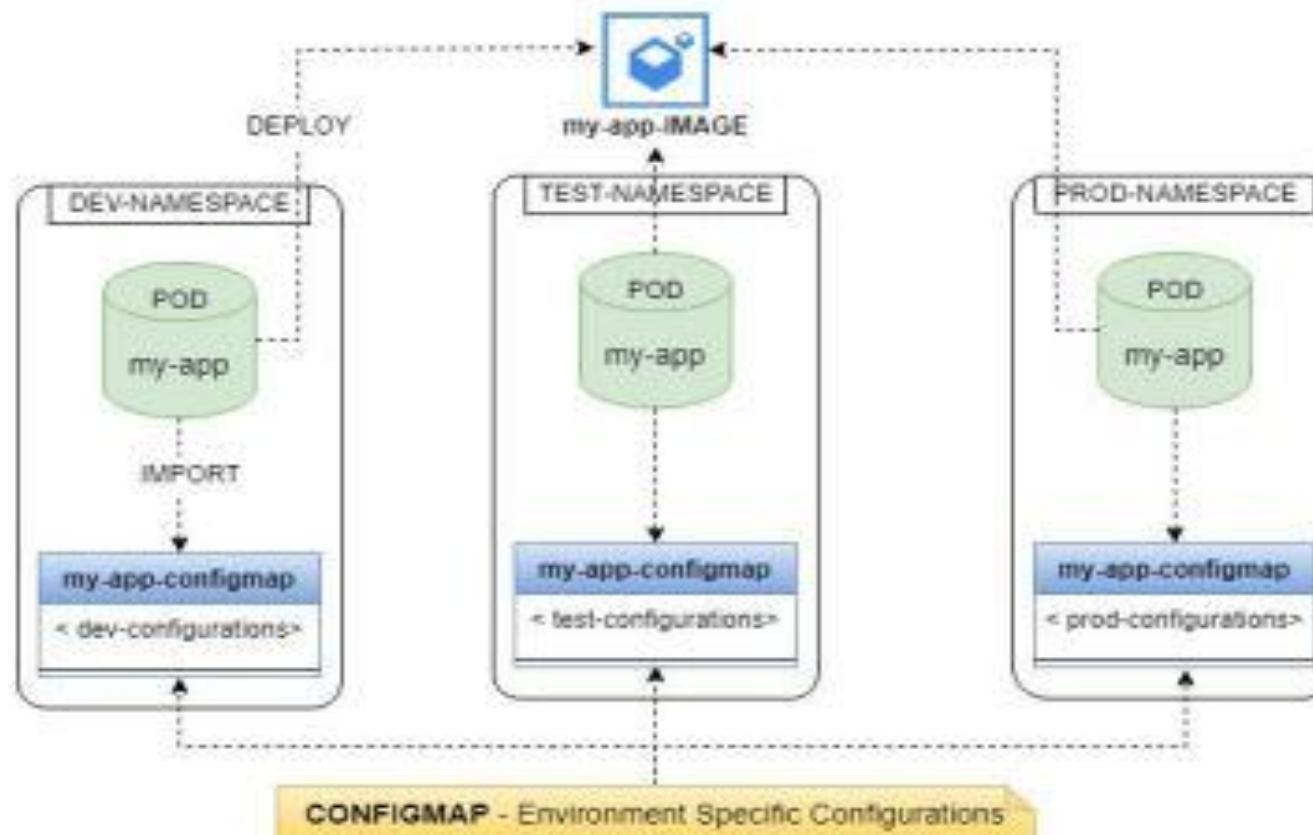
ConfigMap

ConfigMap is an API object used to store non-confidential data in key-value pairs. Pods can consume ConfigMaps as environment variables, command-line arguments, or as configuration files in a volume.

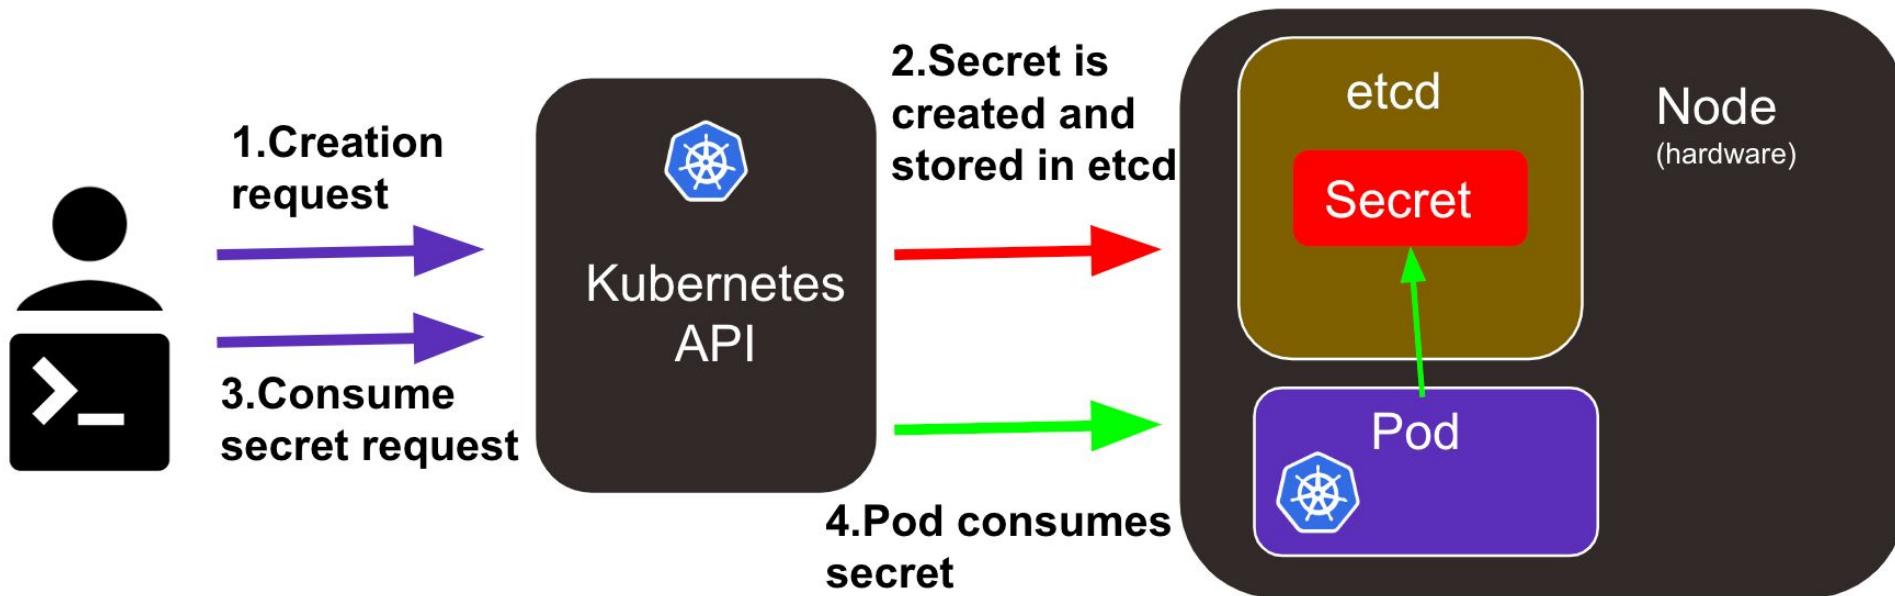


ConfigMap

A **ConfigMap** allows you to decouple environment-specific configuration from your container images, so that your applications are easily portable.



Secrets



A **Secret** is an object that contains a small amount of sensitive data such as a password, a token, or a key. Such information might otherwise be put in a Pod specification or in an image.

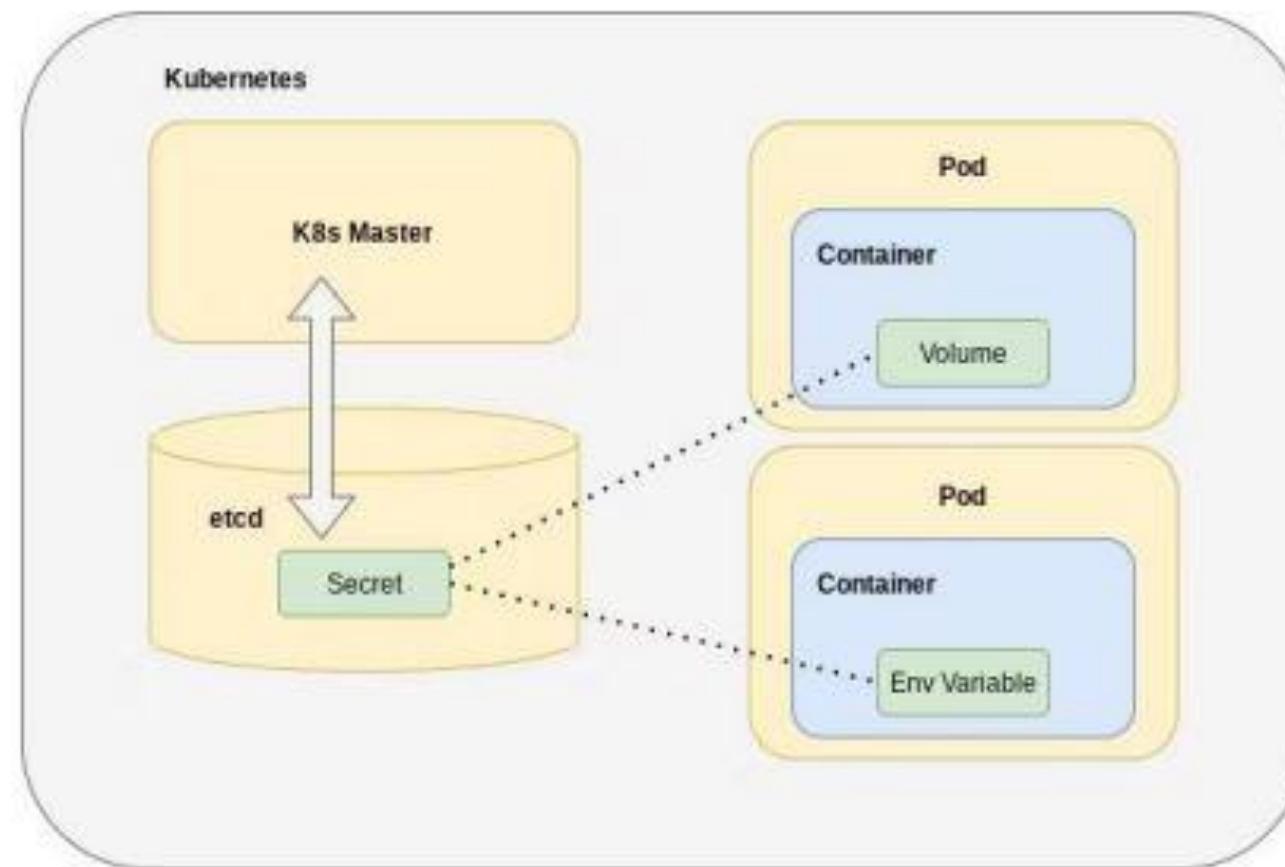
Secrets

Opaque is the default Secret type if omitted from a Secret configuration file. When you create a Secret using kubectl, you will use the generic subcommand to indicate an Opaque Secret type.

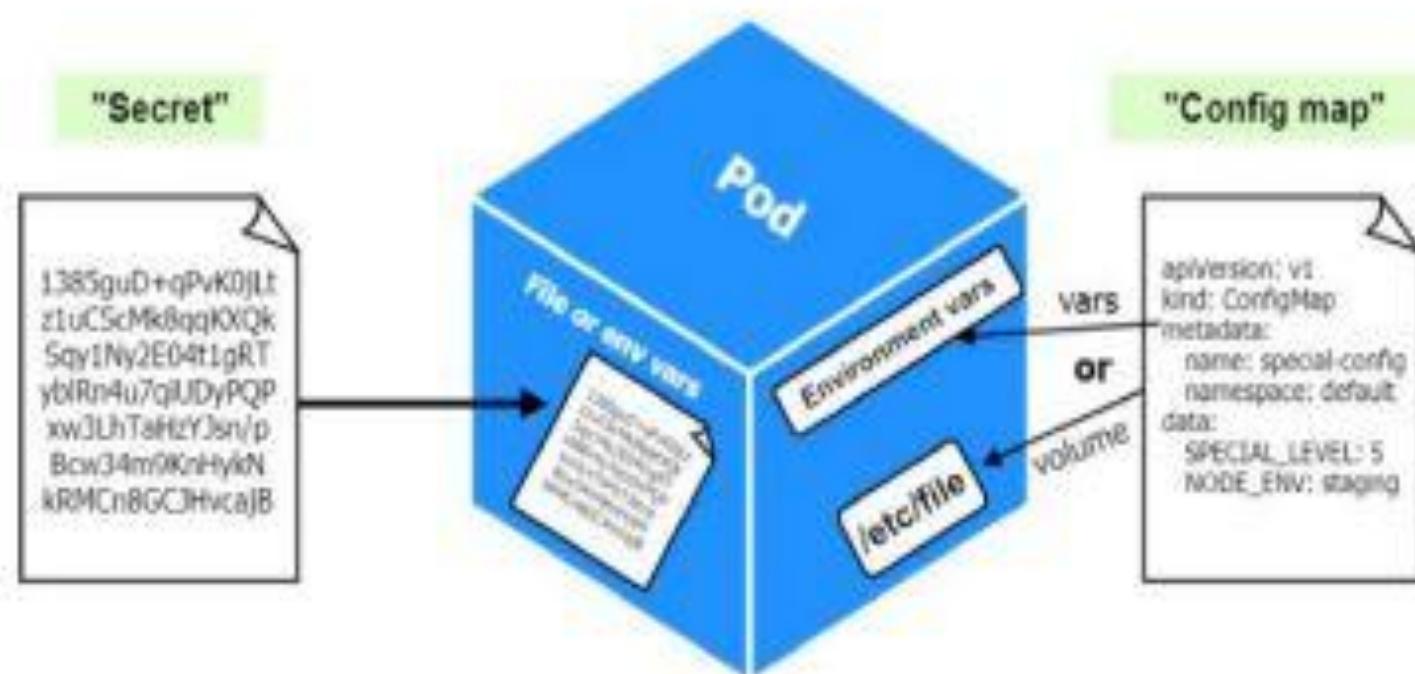
Built-in Type	Usage
Opaque	arbitrary user-defined data
kubernetes.io/service-account-token	ServiceAccount token
kubernetes.io/dockercfg	serialized <code>~/.dockercfg</code> file
kubernetes.io/dockerconfigjson	serialized <code>~/.docker/config.json</code> file
kubernetes.io/basic-auth	credentials for basic authentication
kubernetes.io/ssh-auth	credentials for SSH authentication
kubernetes.io/tls	data for a TLS client or server
bootstrap.kubernetes.io/token	bootstrap token data

Secrets

Secrets can be mounted as data volumes or exposed as environment variables to be used by a container in a Pod.



ConfigMap and Secret





ConfigMap and Secret

```
env:  
  - name: APP_COLOR  
    value: pink
```

```
env:  
  - name: APP_COLOR  
    valueFrom:  
      configMapKeyRef:
```

```
env:  
  - name: APP_COLOR  
    valueFrom:  
      secretKeyRef:
```

1 Plain Key Value

2 ConfigMap

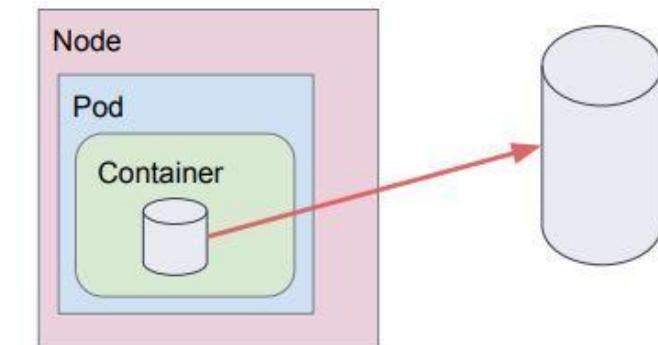
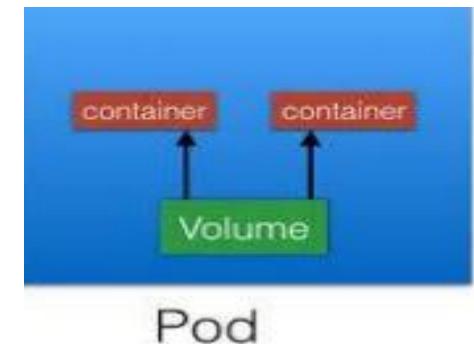
3 Secrets

KUBERNETES Storage

- ▶ Volumes
- ▶ Volume Types
- ▶ PersistentVolume & PersistentVolumeClaim
- ▶ StorageClass

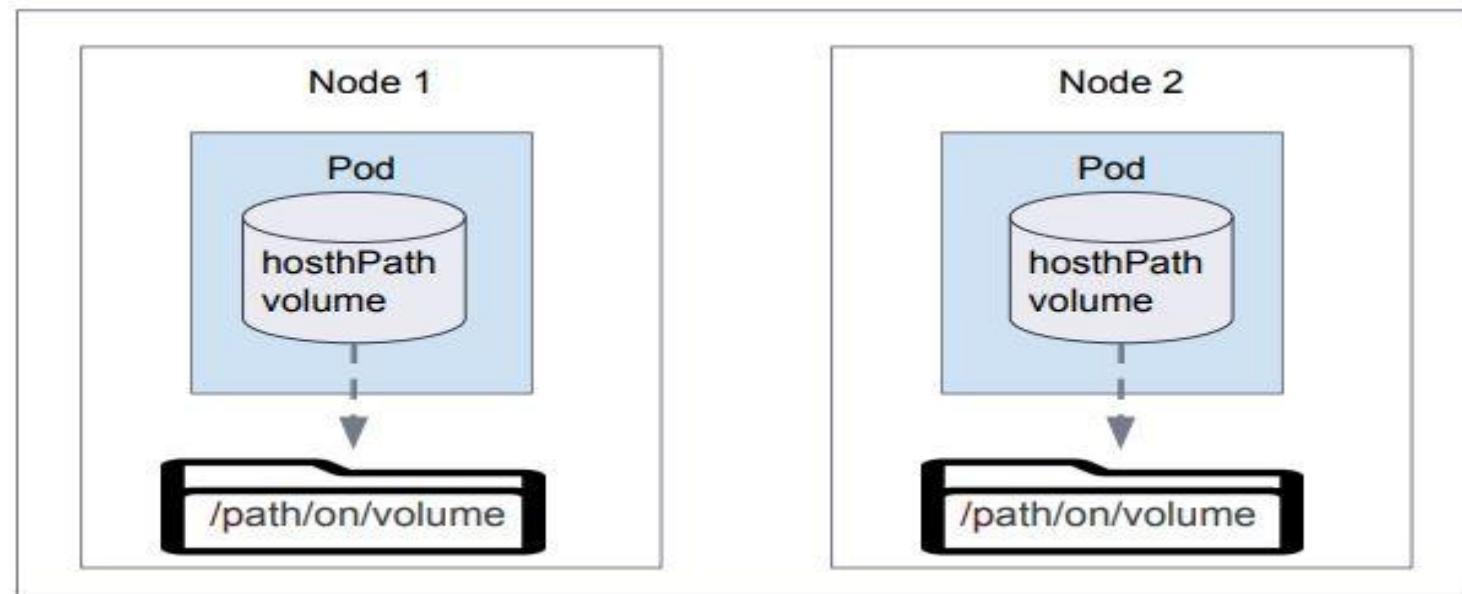
Volumes

- All data stored inside a container is deleted if the container crashes
- When a Container crashes, kubelet will restart it, but the files will be lost which means that it will not have any of the old data.
- To overcome this problem, Kubernetes uses Volumes. A Volume is essentially a directory backed by a storage medium. The storage medium, content and access mode are determined by the Volume Type.



Volume Types

hostPath: A hostPath volume mounts a file or directory from the host node's file system into your Pod.

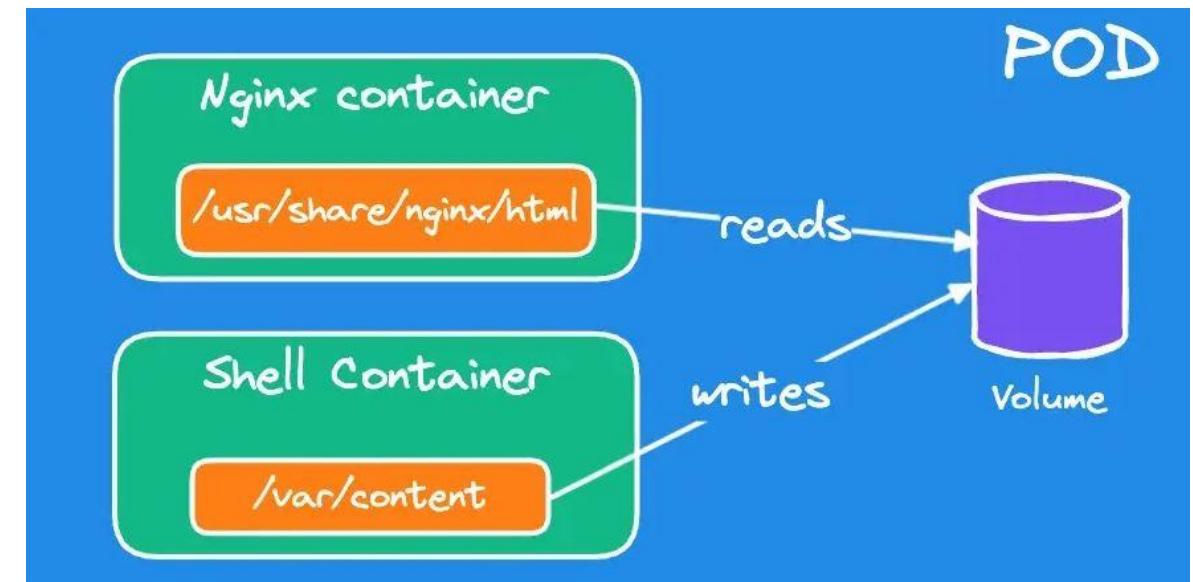


Volume Types

emptyDir: An emptyDir volume is first created when a Pod is assigned to a Node and exists as long as that Pod is running on that node.

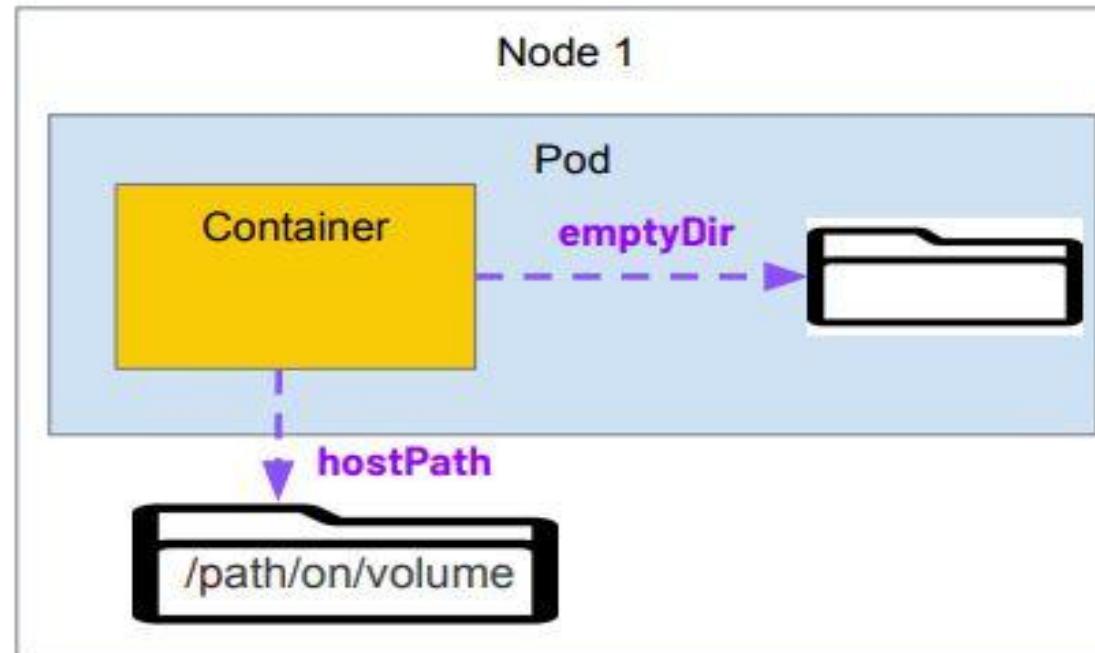
Some uses for an emptyDir are:

- checkpointing a long computation for recovery from crashes
- as a cache



Volume Types

hostPath vs emptyDir



Volume Types

Kubernetes provides several storage plugins that provide access to storage devices deployed in a Kubernetes cluster. These are implemented using the StorageClass object.

Some of the main plugins currently supported by Kubernetes are GCEPersistentDisk, AWSElasticBlockStore, AzureDisk, Glusterfs, NFS, and iSCSI.

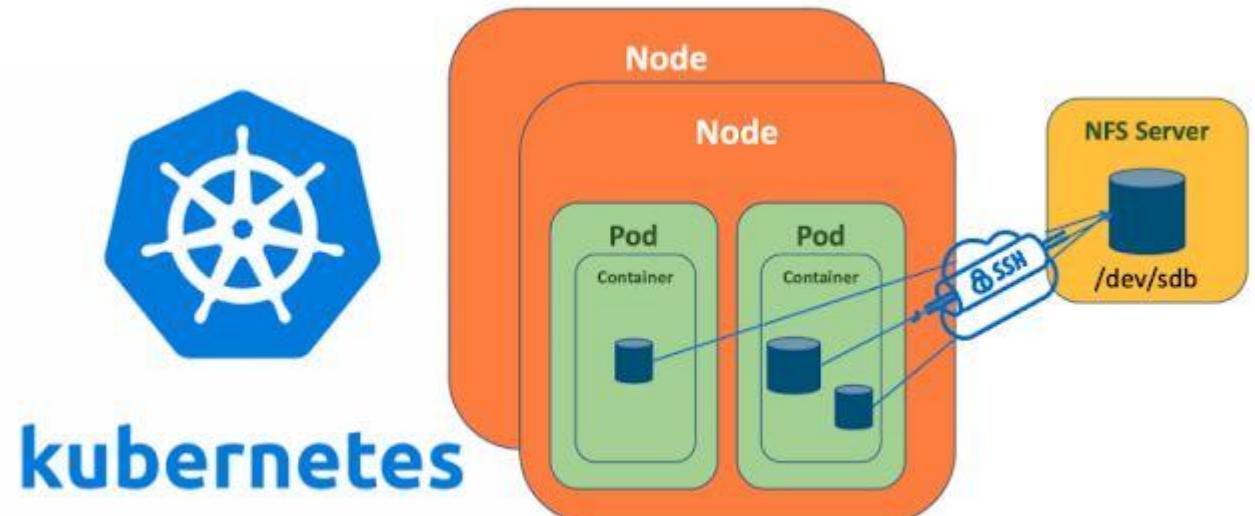
AWS Elastic Block Store: An awsElasticBlockStore volume mounts an Amazon Web Services (AWS) EBS Volume into your Pod.

Azure Disk: An azureDisk is used to mount a Microsoft Azure Data Disk into a Pod.

GCE Persistent Disk: A gcePersistentDisk volume mounts a Google Compute Engine (GCE) persistent disk (PD) into your Pod

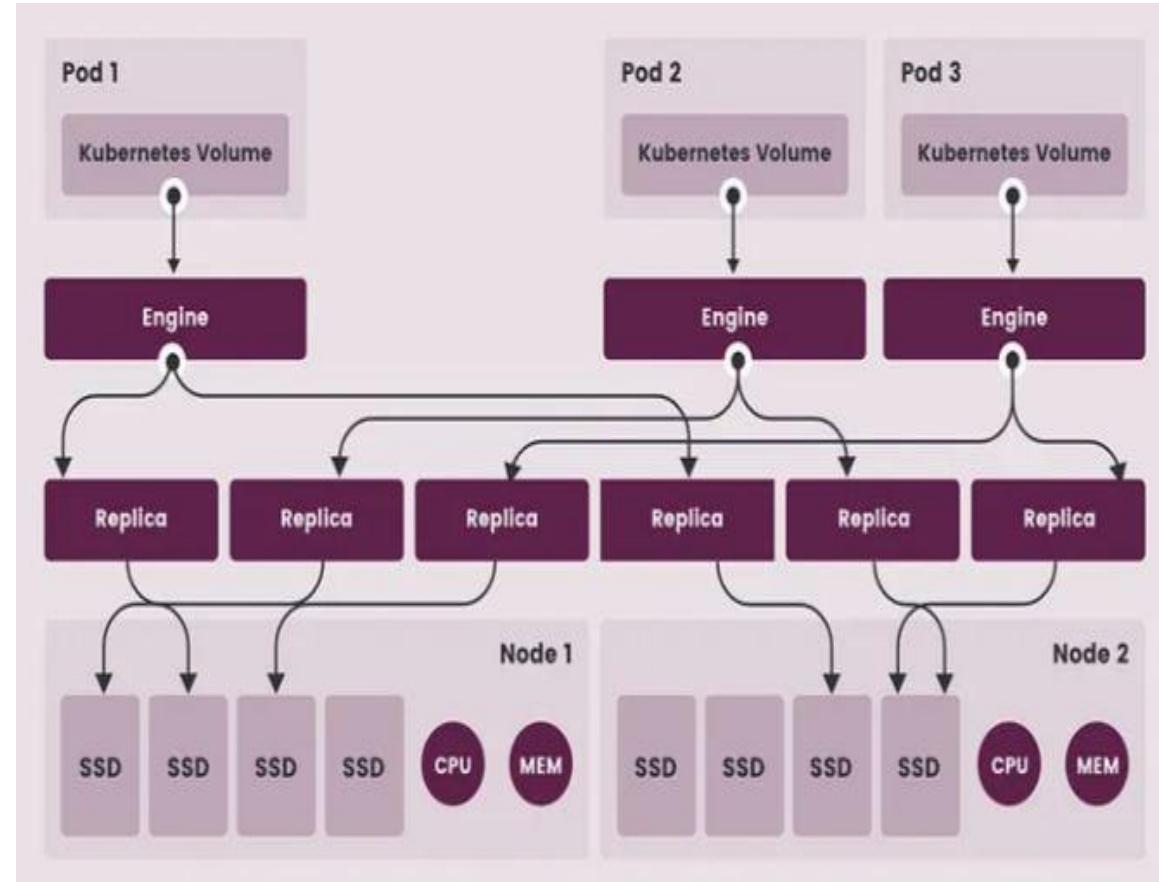
Volume Types

NFS: Network File System (NFS) is a standard protocol for mounting storage devices as local drives. Kubernetes lets you mount an NFS volume as a local drive in a container. Because legacy code often accesses data via NFS, this plugin is very useful for migrating legacy workloads to Kubernetes.



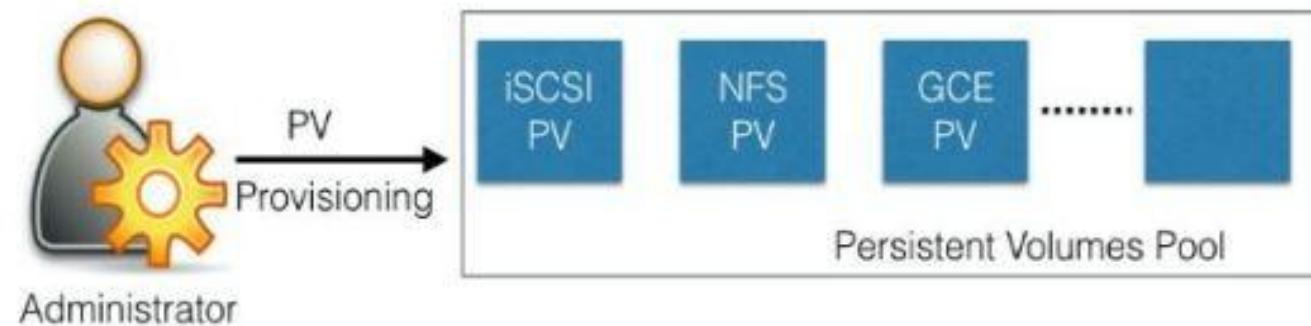
Volume Types

Longhorn is an open-source distributed block storage system designed for Kubernetes. It provides persistent storage for containerized applications running on Kubernetes clusters..Longhorn supports volume snapshots, allowing you to capture the state of a volume at a specific point in time. It also provides backup functionality for disaster recovery purposes.Longhorn uses synchronous replication to ensure that data is stored on multiple nodes, providing high availability. If one node fails, the data is still accessible from replicas on other nodes.



PersistentVolume

A **PersistentVolume (PV)** is a piece of storage in the cluster that has been provisioned by an administrator or dynamically provisioned using Storage Classes



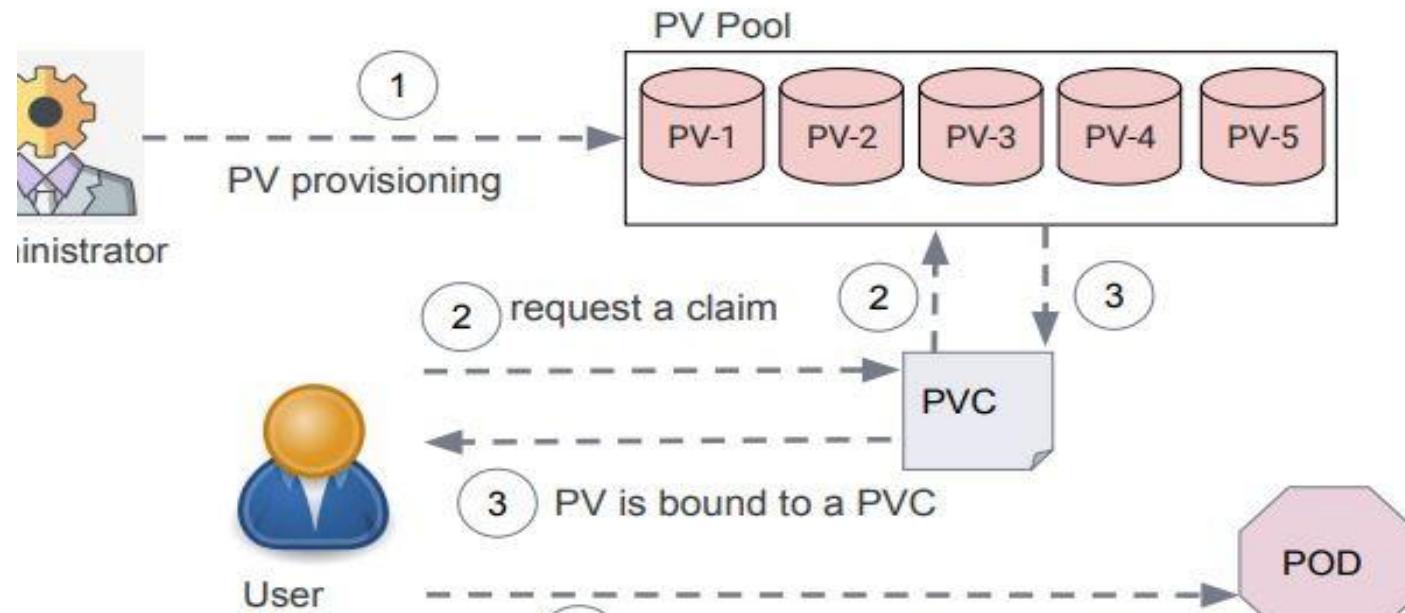
PersistentVolumeClaim

A **PersistentVolumeClaim (PVC)** is a request for storage by a user. Users request for PersistentVolume resources based on type, access mode, and size. There are four access modes:

- ❖ ReadWriteOnce (read-write by a single node)
- ❖ ReadOnlyMany (read-only by many nodes)
- ❖ ReadWriteMany (read-write by many nodes).
- ❖ ReadWriteOncePod (read-write only one pod in the cluster)

Once a suitable **PersistentVolume** is found, it is bound to a **PersistentVolumeClaim**.

PersistentVolumeClaim



StorageClass

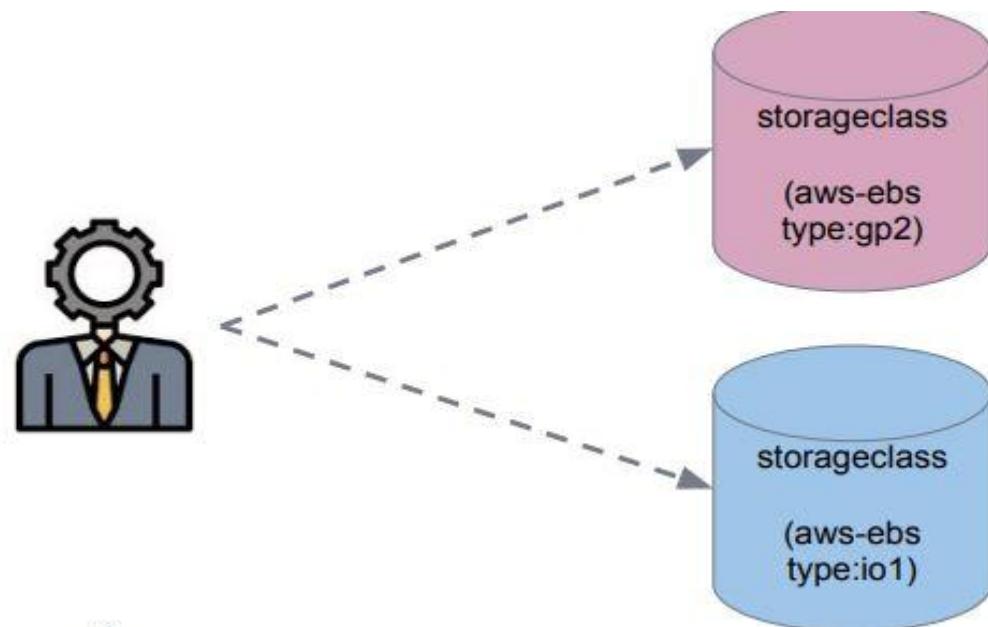
Provisioning: There are two ways PVs may be provisioned; statically or dynamically.

Static: A cluster administrator creates a number of PVs. They carry the details of the real storage, which is available for use by cluster users. They exist in the Kubernetes API and are available for consumption.

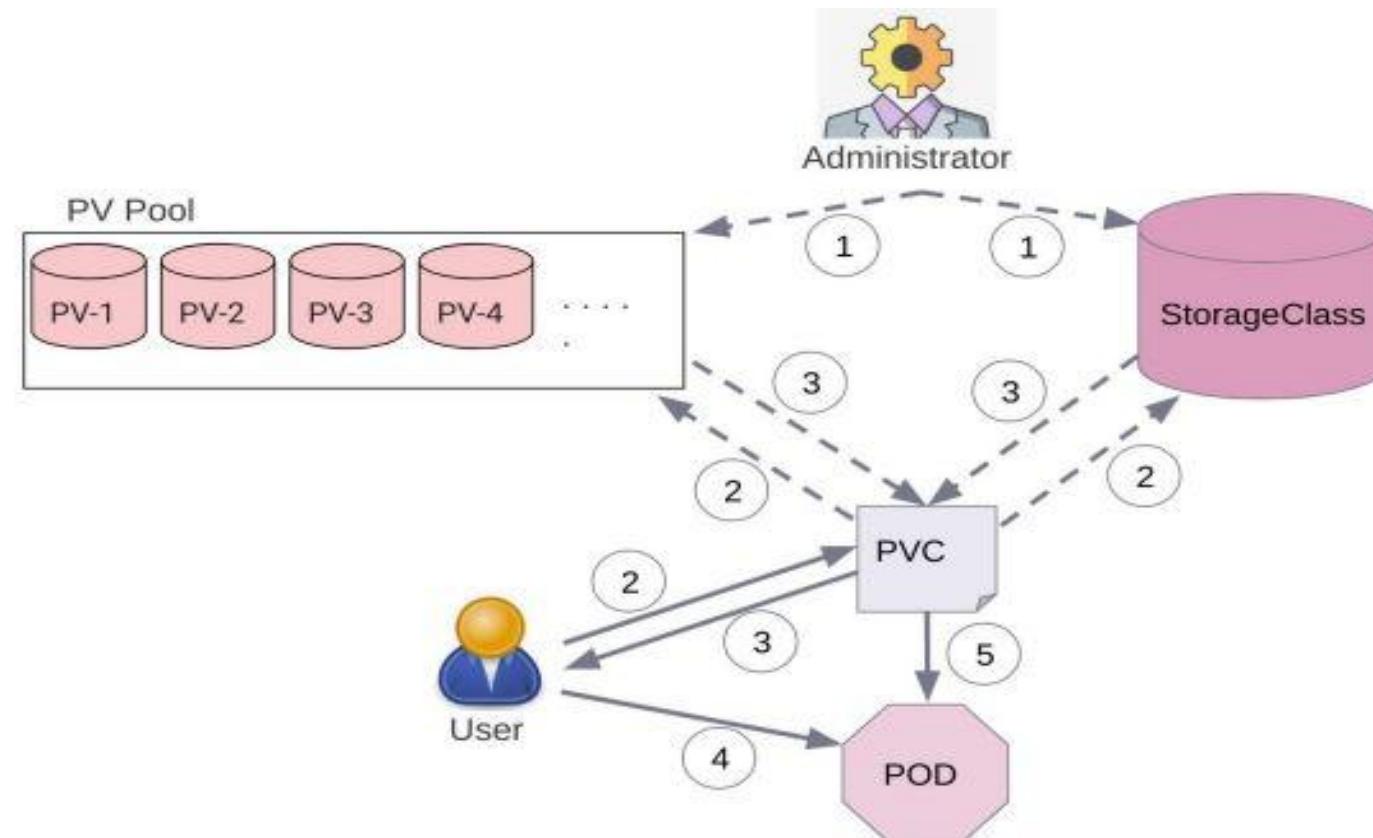
Dynamic: When none of the static PVs the administrator created match a user's PersistentVolumeClaim, the cluster may try to dynamically provision a volume specially for the PVC. This provisioning is based on **StorageClasses**.

StorageClass

StorageClass: provides a way for administrators to describe the "classes" of storage they offer.



StorageClass

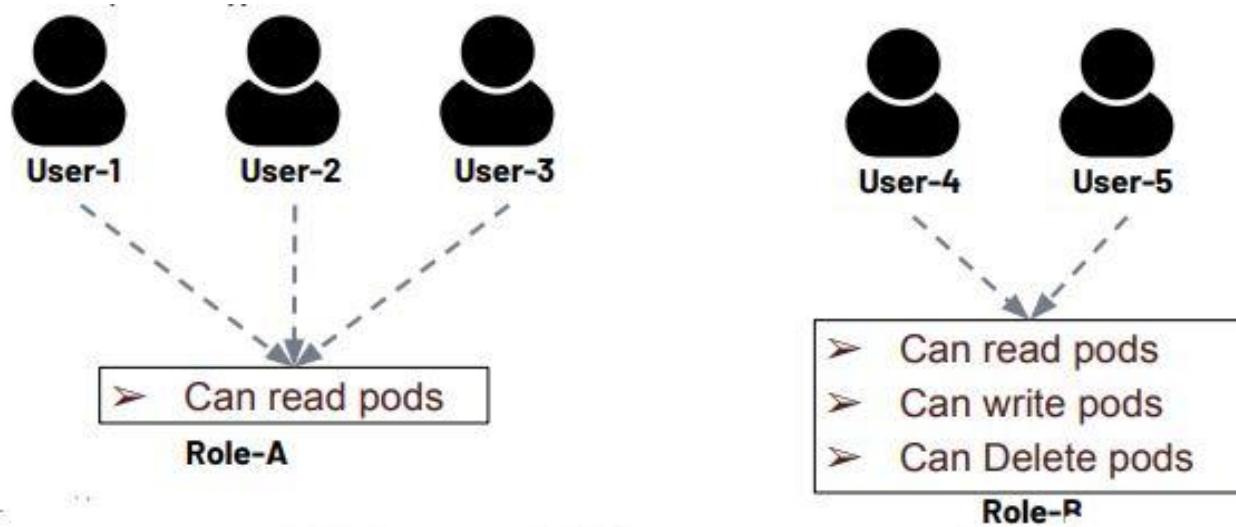


KUBERNETES RBAC

- ▶ Role and ClusterRole
- ▶ RoleBindings and ClusterRoleBindings
- ▶ ServiceAccount

KUBERNETES RBAC

Role-based access control (RBAC) is a method of regulating access to computer or network resources based on the roles of individual users within your organization.



KUBERNETES RBAC

Role and ClusterRole

RBAC Role or ClusterRole contains rules that represent a set of permissions

- **A Role** always sets permissions within a particular namespace; when you create a Role, you have to specify the namespace it belongs in
- **ClusterRole**, by contrast, is a non-namespaced resource

KUBERNETES RBAC

RoleBinding and ClusterRoleBinding

- **A role binding** grants the permissions defined in a role to a user or set of users.
- **A RoleBinding** grants permissions within a specific namespace whereas a ClusterRoleBinding grants that access cluster-wide

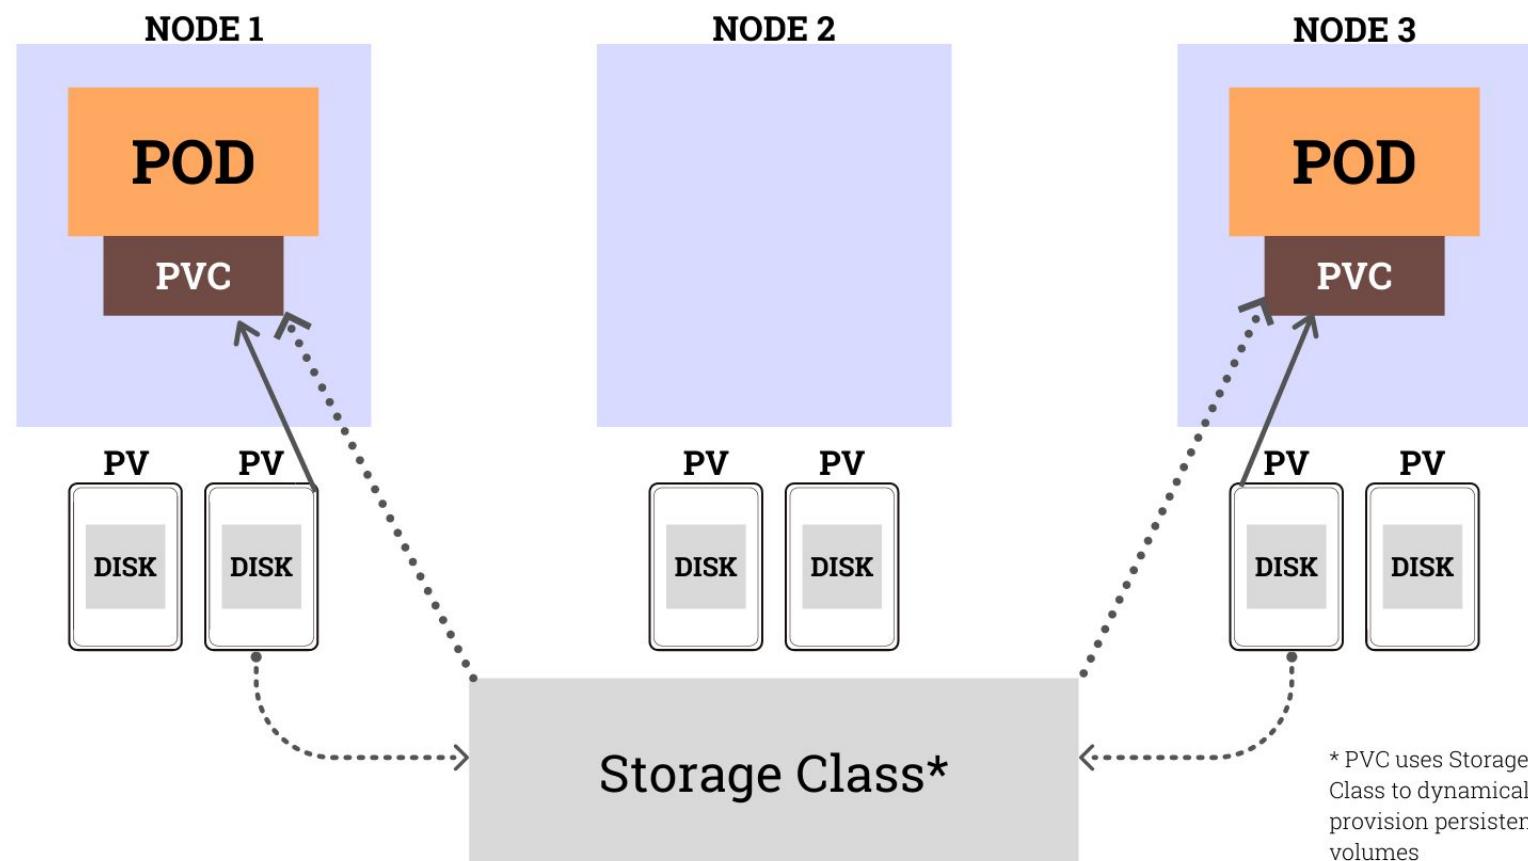
KUBERNETES RBAC

ServiceAccount

A **service account** provides an identity for processes that run in a Pod, and maps to a ServiceAccount object. When you authenticate to the API server, you identify yourself as a particular *user*.

Volumelar ve Depolama

- Kubernetes'te Kalıcı Depolama
- PersistentVolumes (PV): Bağımsız depolama birimleri
- PersistentVolumeClaims (PVC): Depolama kaynakları talep etmek için kullanılır
- StorageClasses:
Sunulan depolama türleri



Helm ve Chartlar

- Helm: Kubernetes Paket Yöneticisi
- Chartlar: Kubernetes uygulamaları için gerekli tüm kaynakları içeren Helm paketleri
- Avantajlar: Kubernetes uygulamalarının dağıtımını ve yönetimini basitleştirir

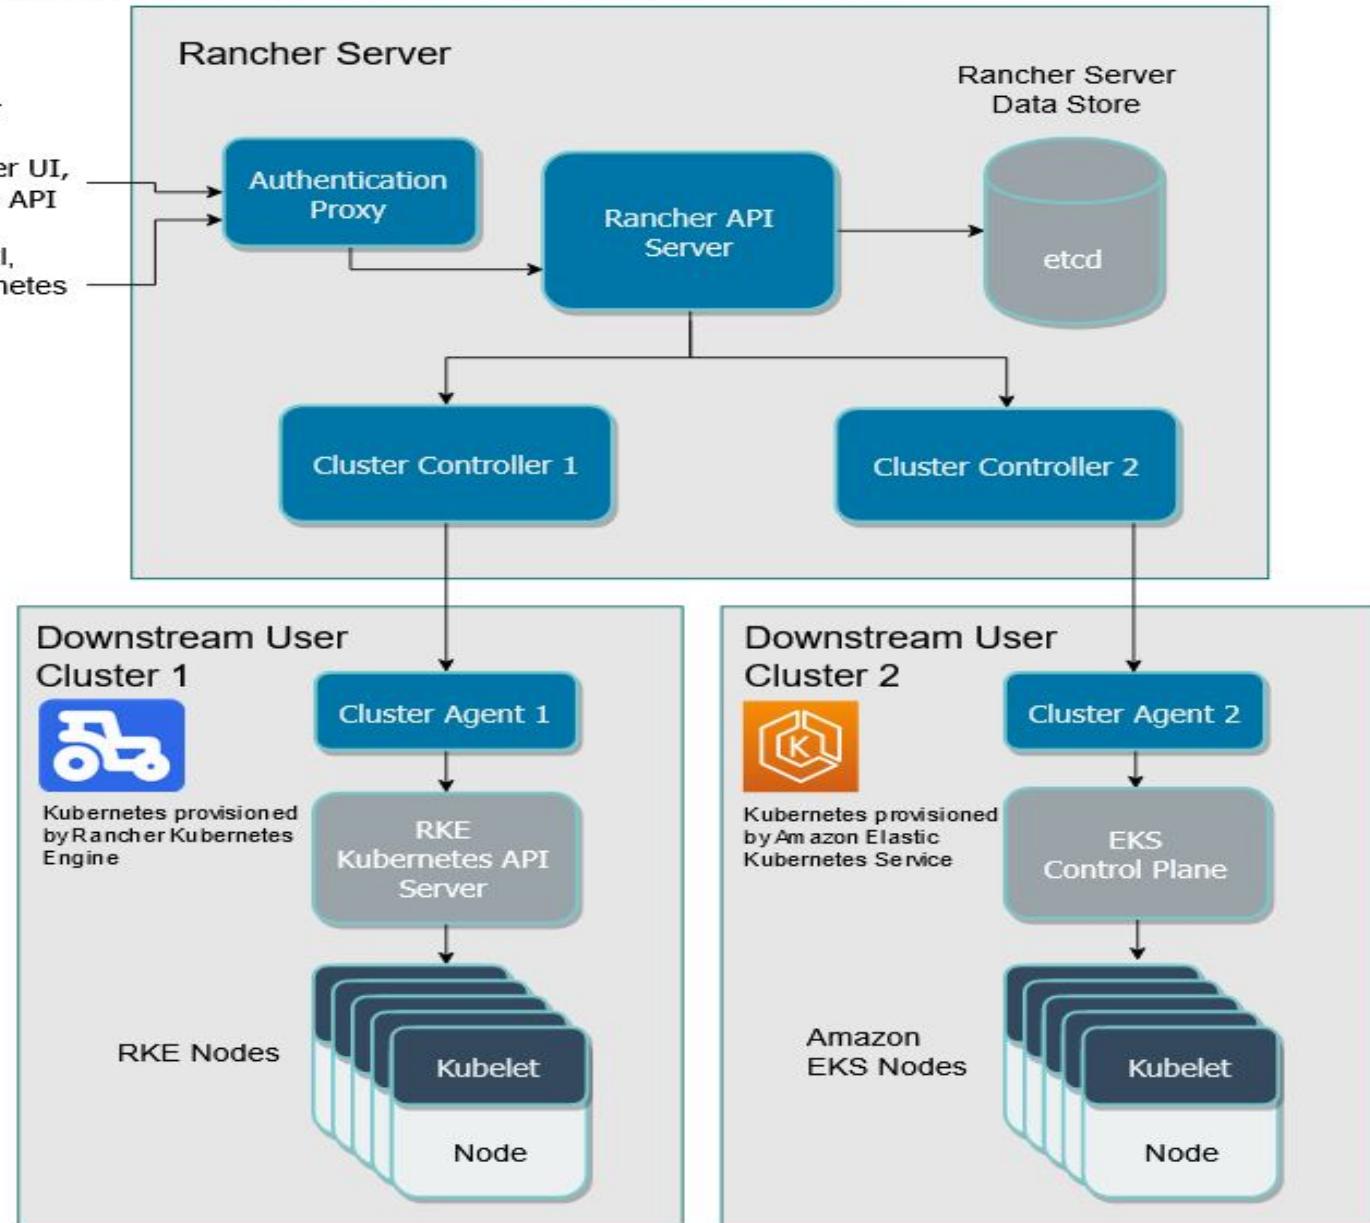
Best Practice'ler

- Cluster Güvenliği: Rol Tabanlı Erişim Kontrolü (RBAC), Gizli Bilgilerin Yönetimi
- Etkili Kaynak Kullanımı: Podların doğru boyutlandırılması, Otomatik ölçeklendirme
- Düzenli Güncellemeler ve İzleme

RANCHER MANAGER MİMARİSİ

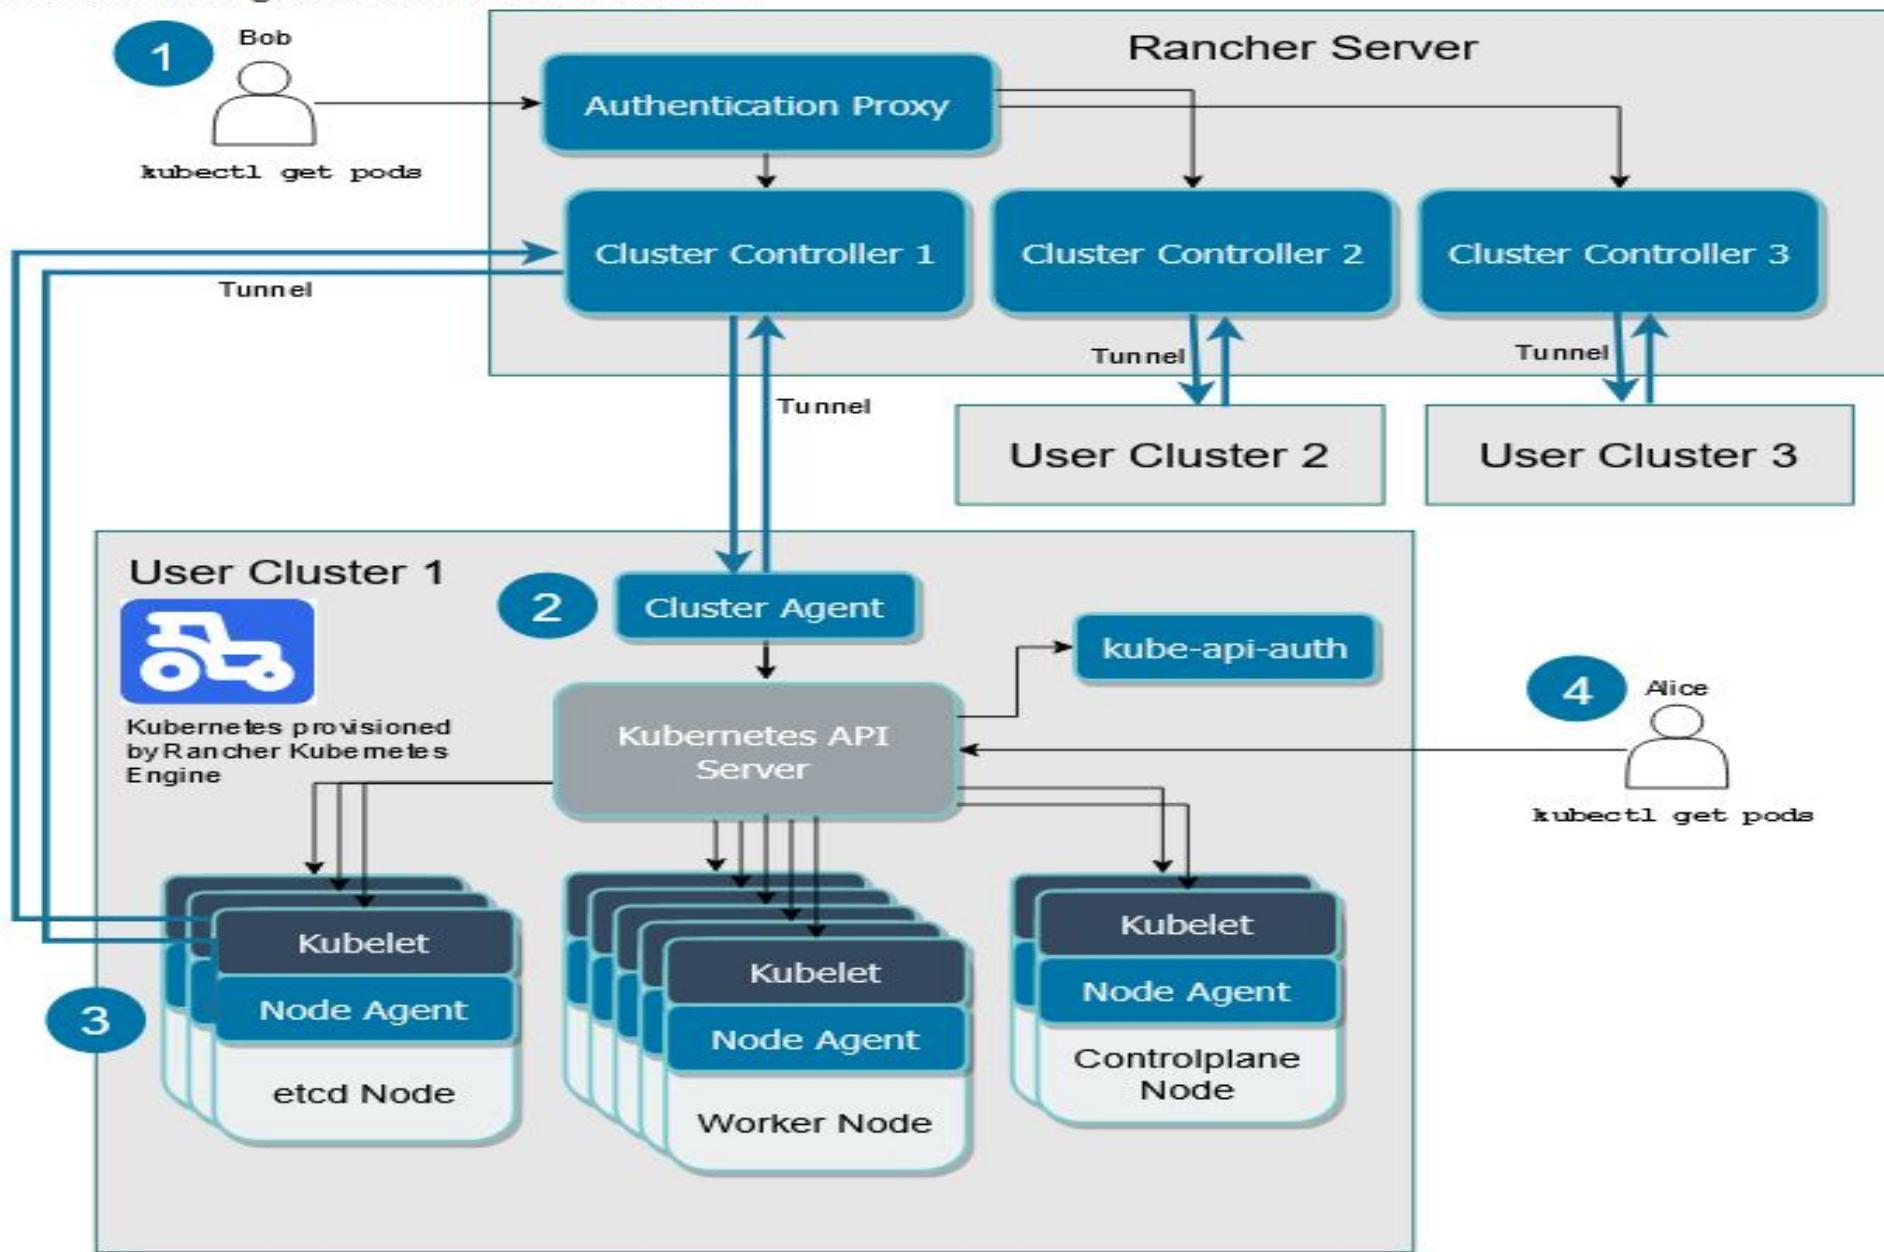


Managing Kubernetes Clusters through Rancher's Authentication Proxy





Communicating with Downstream Clusters

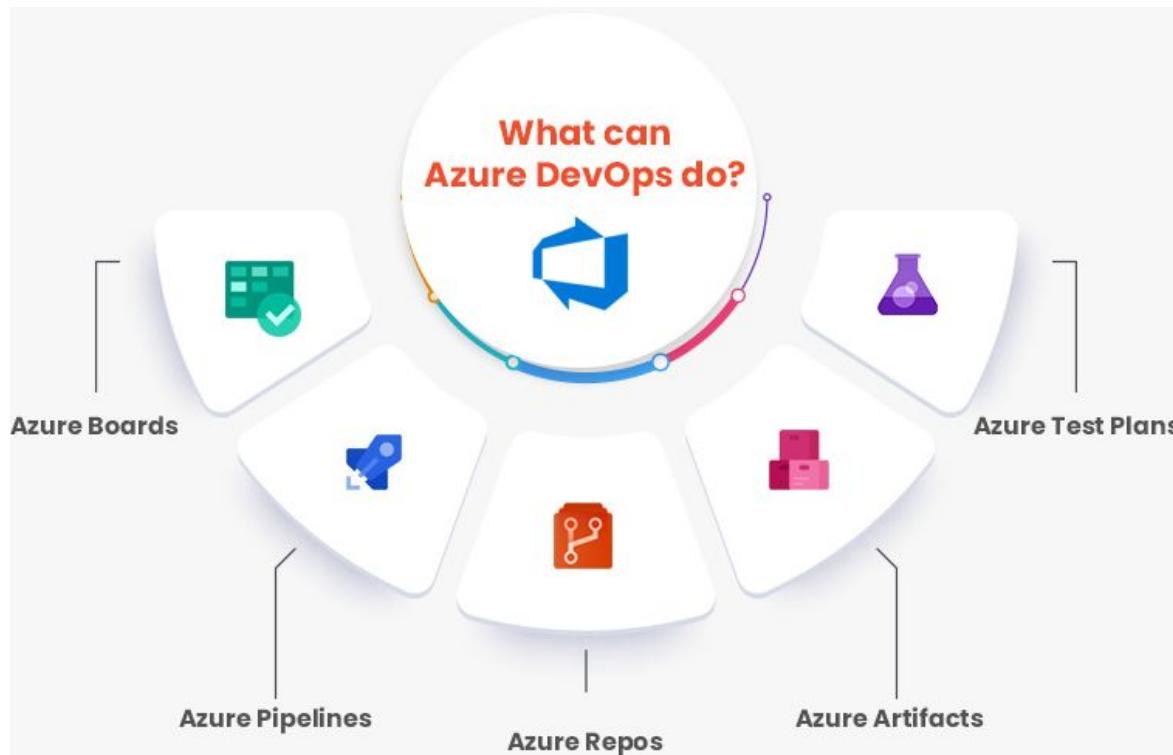


RANCHER MANAGER DEMO

Azure DevOps Server

Azure DevOps'a Giriş

- Azure DevOps Server: CI/CD Süreçlerini Kolaylaştırma
- Bileşenler: Boards, Repos, Pipelines, Test Plans, Artifacts
- Faydaları: İş birliğini artırır, Deployment sıklığını artırır, Ürün kalitesini yükseltir



Azure Repos ve Git Entegrasyonu

- Azure Repos: Sağlam Sürüm Kontrolü
- Git ile Sorunsuz Entegrasyon: Ortak kodlama kolaylaştırır
- Özellikler: Pull Requests, Branch Policies, Repositories Yönetimi

Azure Pipelines: CI/CD

Temelleri

- Azure Pipelines: Otomatik Yapı, Test ve Deployment
- Ana Özellikler: Cross-platform, Popüler diller ve platformlarla entegrasyon
- Baştan sona otomasyon: Koddan buluta

Azure Test Planları ve Test

- Azure Test Planları: Test Etkinliğini Artırma
- Manuel ve Otomatik Testler: Test planlama, yürütme ve takip
- CI/CD ile Entegrasyon: DevOps döngüsünde Sürekli Test

Azure Artifacts ve Paket Yönetimi

- Azure Artifacts: Paket Yönetimini Basitleştirme
- Bağımlılıkları Yönetme: Paketleri barındırma ve paylaşma
- NuGet, npm, Maven ve daha fazlasını destekler

DİNLEDİĞİNİZ İÇİN TEŞEKKÜR EDERİZ



www.hepapi.com



info@hepapi.com



284 Chase Road,
London, N14 6HF,
United Kingdom



Büyükdere Caddesi
No:199, Levent 199,
34330 İstanbul, Turkey